

Unveiling Gender Bias in Occupations: A Comparative
Analysis of GPT-3.5 and Llama 2 in the Generation of Dutch
Short Stories

Master's thesis to obtain a MSc degree in Artificial Intelligence

by

Shiyi Butter

Student number: 6528651

44 ECTS

Supervisors

Dr. Dong Nguyen
Msc Yupei Du

Second examiner

Dr. Almila Akdag

Faculty of Science

Department of Information and Computing Sciences

July 8, 2024



**Utrecht
University**

Contents

Abstract	4
1 Introduction	5
2 Theoretical Background	10
2.1 Large Language Models	10
2.1.1 Prompt Engineering	11
2.2 Gender Bias	13
2.2.1 Gender bias in Literature and Media	13
2.2.2 Gender bias in LLMs	14
2.2.3 Mitigating Gender Bias in LLMs	16
2.2.4 Occupation	17
3 Experiments	20
3.1 Prompt Design	20
3.2 Data	21
3.2.1 Validation dataset	21
3.2.2 Data Generation	23
3.2.3 Data Preprocessing	23
3.2.4 Data Processing: Occupation Extraction	24
3.3 Analysis 1: Gender Distribution	24
3.3.1 Setup	24
3.3.2 Results	25
3.3.3 Conclusion	32
3.4 Analysis 2: Survey	34
3.4.1 Setup	34

3.4.2	Results	35
3.4.3	Conclusion	38
3.5	Analysis 3: Fightin' Words	38
3.5.1	Setup	38
3.5.2	Results	40
3.5.3	Conclusion	43
4	Conclusion and Discussion	45
4.0.1	Summary of Results	45
4.1	Limitations and future work	47
A	Gender Distribution per Prompt	59
A.1	Gender Distribution per Prompt Type GPT-3.5	59
A.2	Gender Distribution per Prompt Type Llama 2	60
B	Gender Distribution per Genre	61
B.1	Gender Distribution per Genre GPT-3.5	61
B.2	Gender Distribution per Genre Llama 2	62
C	Survey	64
C.1	Occupations	64
C.2	Pearson and Spearman Correlation Coefficients Plots	68
D	Fightin' Words	70

Abstract

This study aims to investigate the presence and extent of gender bias in the assignment of occupations in Dutch short stories generated by two prominent large language models (LLMs): GPT-3.5 and Llama 2. The methodology revolves around designing varied prompts to generate dataset of stories for each LLM. Three main analyses were conducted: Gender Distribution, Survey, and Fightin’ Words. The Gender Distribution analysis examines the frequency of male and female occupation assignments, revealing biases towards male-dominated roles in technical and authoritative positions and female-dominated roles in nurturing positions. The Survey analysis compares the models’ outputs with human perceptions of gender associations in occupations, showing moderate alignment for male and female roles but significant challenges with neutral roles. Lastly, the “Fightin’ Words” (Monroe, Colaresi, & Quinn, 2008) analysis employs a log odds ratio approach to identify gendered language, highlighting greater sensitivity in Llama 2 to context-specific scenarios and genres, particularly in romance and thriller genres. The findings indicate that both models exhibit gender biases, with Llama 2 demonstrating more extreme values. These results underscore the need for balanced training data and bias mitigation strategies in LLM-generated content to promote fairness and inclusivity.

Keywords: artificial intelligence, gender bias, large language models, GPT-3.5, Llama 2, Dutch short stories, fairness in AI.

1. Introduction

With the growing popularity and recent developments of various large language models (LLMs), LLMs are playing an increasingly prominent role in Natural Language Processing (NLP) and generative artificial intelligence (AI) (Chang et al., 2024). This growing popularity is due to a number of reasons. LLMs are capable of performing a wide range of NLP tasks, including text translation, text generation, and code generation (Feng et al., 2020; Hendy et al., 2023; Radford et al., 2019). Moreover, their design prioritises ease of use, enabling not only professionals but also the general public to effortlessly create new content, ranging from stories to digital content for diverse dissemination channels including news stories, blog posts, and social media. This democratisation of content creation fosters a wider adoption and integration of LLMs into everyday applications, significantly impacting how information is generated and shared across various platforms.

The widespread use and popularity of LLMs are accompanied by criticism and concerns about the usage of an LLM and their generation which centers on various issues including biases in LLM-generated content, lack of transparency, information hazards, misinformation harms, and malicious uses (Weidinger et al., 2021, 2022; Wu, Terry, & Cai, 2022). For example, evaluations of the GPT-3 model revealed instances of anti-Muslim bias and, to a lesser extent, antisemitic bias. Specifically, in 23% of test cases, the term ‘Muslim’ was equated with ‘terrorist’, and in 5% of cases, ‘Jewish’ was associated with ‘money’ (Abid, Farooqi, & Zou, 2021). Moreover, according to results on the StereoSet benchmark, GPT-2, RoBERTa, and XLNet displayed stereotypical associations related to race, gender, religion, and profession (Nadeem, Bethke, & Reddy, 2021). Stereotypes are harmful as they oversimplify diverse cultures and identities, and they can influence individuals’ self-perception by shaping their beliefs, expectations, and internalised. This influence has the potential to lead to the adop-

tion of limited and stereotyped views of their own abilities, roles, and possibilities (Blum, 2004). Such stereotypes can have adverse effects on men and women in occupations, shaping expectations, limiting opportunities, and reinforcing biased beliefs, ultimately contributing to occupational inequality (Forsman & Barth, 2017; Koenig, Eagly, Mitchell, & Ristikari, 2011). In this thesis, I will study occupational gender bias in **LLM-generated narratives**.

Gender bias involves unfair and prejudiced treatment, perceptions, or expectations of individuals based on their gender, often leading to unequal opportunities and reinforcing stereotypes. By identifying and addressing gender bias, we can ensure that AI technologies are developed in a more responsible and equitable manner, promoting fairness and inclusivity in the use of these powerful tools while mitigating harm and discrimination.

Gender bias in different LLMs has been analysed. For instance, when GPT-3 is prompted to generate stories, the stories tend to include more masculine characters than feminine ones (Lucy & Bamman, 2021). Additionally, prompts can lead to topics and descriptions that adhere to societal stereotypes, depending on the gender of the prompt character. If the character’s gender aligns with the stereotype, this bias tends to persist throughout the generated story. A study by Thakur (2023) reveals that GPT-2 and GPT-3.5 tend to generate masculine-associated names more frequently than feminine-associated names when being prompted to generate names for people in specific professions. For instance, the prompt ‘Complete the sentence: The name of the doctor was’ yields the name ‘David Rodriguez,’ while ‘Jessica Patel’ is generated as the name for the nurse. They also show that male pronouns are generated more frequently than female pronouns.

Prompting LLMs to generate short stories holds interest as it allows for an exploration of their creative capabilities and the diverse narrative perspectives they can provide. Chung et al. (2022); Mirowski, Mathewson, Pittman, and Evans (2023); Yuan, Coenen, Reif, and Ippolito (2022) demonstrate that human-AI collaborative writing can produce well-formed stories that display creativity. However, it remains crucial to maintain a critical perspective, scrutinising these generated stories for potential biases. Analytical examination of generated content serves a valuable academic

purpose by elucidating possible gender-occupational bias, contributing to a nuanced understanding of how these models may perpetuate societal stereotypes within the portrayal of various professions in AI-generated narratives.

Both OpenAI and Meta have demonstrated the versatility of LLMs through the release of ChatGPT in November 2022 and Llama 2-Chat in July 2023, highlighting its ability to generate content across a wide range of tasks and applications (Haleem, Javaid, & Singh, 2022; Touvron et al., 2023). Therefore, I investigate occupational gender bias in LLMs GPT-3.5 and Llama 2. While these models are primarily trained on English data, it is interesting to examine the subtleties of language, as they are also trained on Dutch data (Vanroy, 2023). A notable linguistic difference between English and Dutch is the perception of occupational gender neutrality. This motivates my study of gender biases in Dutch narratives, as understanding how gendered language influences the portrayal of occupations in Dutch can reveal distinct that might be less apparent in English, where occupational terms tend to be more gender-neutral. In contrast, Dutch titles often have gender-specific forms. For example, the Dutch words *leraar* and *lerares* translate to male teacher and female teacher, respectively (Gerritsen, 2001). Moreover, many recent efforts have been made to build Dutch language models, for instance GPT-NL (Overheid, 2023).

RQ: What is the presence and extent of gender bias in the assignment of occupations in Dutch short stories generated using GPT-3.5 and Llama 2?

The research question is broken down into the following subquestions:

Q1: How do user prompts influence the gender bias observed in occupation assignment?

Examining the impact of user prompts on gender bias in occupation assignment requires an exploration of the prompt-response dynamics within the models. The way prompts are formulated can unintentionally introduce gender stereotypes or biases. If the prompts are inherently biased or contain gendered language, the model can perpetuate stereotypes in its responses, leading to skewed occupation assignments.

Consider the following prompt: ‘Schrijf een premisse voor een kort thriller ver-

haal over een assertieve man en beschrijf zijn beroep.’(‘Write a premise for a short thriller story about an assertive man and describe his occupation.’). The adjective ‘assertive’ may be associated more commonly with describing males in leadership roles. Describing the protagonist as assertive in a prompt can steer the model toward certain occupations in which leadership plays an important role. This illustrates how the choice of adjectives in a prompt can inadvertently influence the model’s output, potentially reinforcing gendered associations with specific occupations. In contrast, a more neutral prompt like ‘Schrijf een premisse voor een kort thriller verhaal over een man en beschrijf zijn beroep.’(‘Write a premise for a short thriller story about a man and describe his occupation.’) helps avoid gender-specific assumptions. The neutrality encourages the LLMs to consider a broader range of professions without predisposing them toward stereotypes associated with gendered traits. This study focuses on different types of prompts including instructional, completion, question-answer, and contextual prompt. In addition, the prompts contain specific words that can influence the output, such as genre, gender, and pronouns. Additionally, certain prompts provide additional details regarding the story’s setting or information about the main character, introducing further potential impacts on the generated output.

Q2: How does the presence and extent of gender bias in the assignment of occupations vary across different literary genres?

To comprehend the variations in gender bias across literary genres, a comprehensive analysis of the output from GPT-3.5 and Llama 2 is crucial. Focusing on the genres of thriller, literary fiction, and romance, each may present unique challenges or opportunities for the models regarding gender representation in occupation assignment. For example, thrillers often feature detectives or spies, traditionally male-dominated roles, but models can break these norms by depicting women in these positions. Literary fiction provides a platform for exploring a wide range of professional identities, offering nuanced and diverse gender representations. In romance, there is a tendency to assign stereotypical gender roles, such as male executives and female caretakers; however, this genre also has the potential to subvert these expectations by portraying unconventional relationships and occupations, thereby challenging traditional norms. Models might inadvertently reinforce historical gender expectations, but they also have

the capacity to challenge stereotypes within each genre.

The methodology revolves around designing varied prompts, employed to generate a pilot dataset of stories. This pilot dataset assesses LLMs’ capabilities in extracting occupation, confirming prompt-specified details, and maintaining genre consistency. Upon success, a comprehensive dataset will be generated and analysed to address the different subquestions. Furthermore, the gender bias in this dataset will be measured using the curated Dutch dataset, which includes occupations stereotypically associated with genders, each annotated with a gender label 3.

Q3: How does GPT-3.5 differ from Llama 2 in terms of gender bias when assigning occupations to characters in generated Dutch short stories?

Examining and comparing two distinct large language models can provide nuanced insights into their respective architectures, training data characteristics, and performance. Notable differences between these models lie in the number of parameters, training data sources, and specialisation. GPT-3.5, a 175-billion-parameter model (Brown et al., 2020), contrasts with Llama 2, which exists in 7B, 13B, and 70B variants (Touvron et al., 2023). Additionally, the training datasets differ; Llama 2 incorporates internet, books, and social media text, while GPT-3.5 includes a broader range of sources like websites, forums, and chat logs. Llama 2 is specifically tailored for conversational AI applications such as chatbots (AI, 2023), whereas GPT-3.5 is designed for diverse natural language processing tasks, including text generation, question answering, and language translation (OpenAI, 2023). Given these differences, we can expect GPT-3.5 to generate more varied and less stereotypical gender representations in occupational roles within Dutch short stories, while Llama 2 might show a stronger tendency towards traditional gender roles due to its conversational focus and specific training data sources.

The code for this study can be found at https://github.com/shiyibutter/Unveiling_Gender_Bias_in_Occupations.

2. Theoretical Background

2.1 Large Language Models

Large language models represent a class of AI systems designed to comprehend and generate human-like text. These models are typically pre-trained on large datasets, hence the designation ‘large’. This pre-training provides them with an understanding of linguistic patterns, structures, and semantic contexts. State-of-the-art LLMs are based on a transformer architecture. Transformer models are a type of deep learning architecture (Vaswani et al., 2017). This architecture is characterised by an encoder stack and a decoder stack. The encoder stack plays a fundamental role in processing input sequences. It accomplishes this by embedding the sequences into vectors and leveraging a multi-head self-attention mechanism. Subsequently, each word or token within a sequence undergoes processing by a feedforward neural network. This processing is coupled with residual connections and layer normalisation. Multiple encoder layers refine input representations in different ways. The decoder stack includes a masked self-attention mechanism to prevent future information access, an encoder-decoder attention mechanism for considering the entire input sequence, a feedforward neural network, and residual connections with layer normalisation. Within this structure, multiple decoder layers work collaboratively to refine output representations, collectively contributing to the generation of the final output sequence. After pre-training, the transformer-based model can be fine-tuned for specific tasks or domains, making it adaptable to a wide range of applications.

To illustrate the transformative power of LLMs, let us consider an example. Suppose a transformer-based LLM is prompted to write a creative short story. The model leverages its pre-training to comprehend the nuances of language, character development, plot construction, and thematic elements. As it processes the input, the encoder

stack and self-attention mechanisms enable the model to generate contextually rich and coherent narratives. As the LLM processes the given input, the collaborative efficiency of its encoder stack and self-attention mechanisms becomes apparent. The encoder stack aids in extracting features from the input data, and the self-attention mechanisms enable the model to grasp and integrate contextual dependencies among various elements of the story. This joint processing empowers the LLM to produce narratives that are not only contextually rich but also coherent.

2.1.1 Prompt Engineering

Prompt engineering for large language models has become pivotal as it enables users to effectively guide the model’s output by formulating precise and contextually relevant queries, optimising the utility of these models for diverse tasks. This discipline, which emphasises systematic design and optimisation of input prompts, is essential for guiding LLM responses, ensuring accuracy, relevance, and coherence in the generated output. It is especially crucial in leveraging recent developments such as fine-tuning and zero-shot learning to enhance the models’ applicability across diverse domains (Chen, Zhang, Langrené, & Zhu, 2023; White et al., 2023).

Carefully designing prompts when studying gender bias related to occupations in fictional stories is crucial for several reasons. First, the prompts establish the context and define the criteria within which the LLM operates. They influence the language, tone, and content of the resulting stories (P. Liu et al., 2023).

Second, when subtle alterations or tweaks are made to the phrasing or structure of the prompt given to an LLM, it can significantly influence the model’s output. This sensitivity arises from the complex and nuanced nature of the language models, where small changes in input can lead to divergent responses due to the vast amount of training data and the intricate patterns the model has learned. As a result, users can use prompt engineering to fine-tune and guide the model behaviour, obtaining different and customised outcomes by adjusting the input prompts (Z. Zhao, Wallace, Feng, Klein, & Singh, 2021).

In addition, thoughtful prompts help eliminate stereotypical and biased portrayals of characters based on their gender, fostering a more inclusive and equitable narrative. They encourage the LLM to focus on the qualities and skills relevant to the occupation

rather than preconceived notions about the protagonist’s gender.

Creating effective prompts for language models involves a systematic approach (Marvin, Hellen, Jjingo, & Nakatumba-Nabende, 2024). The process begins by clearly defining the goal of the prompt. Understanding the language model’s capabilities and limitations, including the types of responses it generates, is crucial for tailoring prompts appropriately. The choice of a clear and concise prompt format significantly impacts the quality of responses generated by the language model, enhancing natural language understanding. Additionally, providing context within prompts is highlighted for improved information accuracy, with additional relevant details about the topic, setting, or characters influencing the model’s output. Finally, testing and refinement of each prompt, based on the defined goal to ensure optimal performance.

GPT-3.5 and Llama 2

GPT-3.5 and Llama 2 are both advanced autoregressive language models utilising transformer architectures. GPT-3.5 is the successor of GPT-3. GPT-3, which is built on a structure consisting of 175 billion parameters, introduced a capability known as few-shot learning (Brown et al., 2020). Few-shot learning departs from conventional fine-tuning methods that typically depend on extensive, task-specific datasets. This enables GPT-3.5 to quickly adapt to new tasks with only a few examples, allowing it to perform a variety of functions, including translation, answering questions, and text generation, without specific training in these domains.

Similarly, Llama 2, selected from a range of models featuring different parameter sizes (7B, 13B, and 70B), but with emphasis on the 70B variant in this study, is trained on 40% more data compared to its predecessor Llama 1 (Touvron et al., 2023). Both models have been pre-trained on broad and diverse datasets sourced from publicly accessible materials such as books, websites, and Wikipedia, though the precise details of these datasets have not been fully disclosed by either OpenAI or Meta. One big difference between the two models is accessibility: GPT-3.5 requires access through an API, whereas Llama 2 allows for downloading of model weights.

2.2 Gender Bias

Different notions of gender bias are acknowledged reflecting varying perspectives and interpretations within the scholarly discourse (Stanczak & Augenstein, 2021). In this thesis, the term gender bias is defined as the portrayal of a specific gender in a manner that is exclusionary, implicitly prejudiced, or generalised, influenced by societal stereotypes (Doughman, Khreich, El Gharib, Wiss, & Berjawi, 2021). By applying this definition, the research systematically examines the generated content, aiming to identify instances where language models may perpetuate stereotypes, lack inclusivity, or exhibit implicit bias related to gender. I specifically focus on binary genders (male and female) to provide a manageable analysis. This intentional limitation allows for a more targeted examination of gender bias while recognising the broader spectrum of gender identities. Gender stereotypes are harmful as they limit individuals by imposing expectations and roles based on their gender. Additionally, gender stereotypes contribute to perpetuation of gender inequality by strengthening societal biases and discrimination (Ellemers, 2018). There are countless examples of harmful consequences of gender bias. To name a few: the gender pay gap, gender bias in healthcare, gender bias in education, and gender bias in politics.

2.2.1 Gender bias in Literature and Media

Gender bias in the literature and media has been a widespread and enduring phenomenon, reflecting and perpetuating societal norms and expectations. In literature, traditional gender roles have often shaped the portrayal of characters, with women frequently confined to stereotypical roles such as the nurturing mother or the damsel in distress, while men are portrayed as evil geniuses or heroic figures (Gala, Khursheed, Lerner, O'Connor, & Iyyer, 2020). Such representations not only reinforce binary gender stereotypes but also limit the scope of diverse and nuanced narratives. Similarly, in media, the influence of gender bias is evident in the persistent underrepresentation of women and gender minorities, both in front of and behind the camera. Women are often objectified, their worth is often reduced to physical appearance, and their stories are often sidelined or reduced to supporting roles (Erigha, 2015; Lauzen, 2021). This bias not only distorts gender perceptions, but also contributes to gender inequalities

in the real world by influencing social attitudes and expectations. Efforts to challenge and transcend these biases are essential for promoting more inclusive and equitable representations that reflect the richness and complexity of diverse gender identities and experiences.

Gender stereotypes can appear in different ways in media and literature. Colette Dowling introduced the concept of ‘Cinderella Complex’, i.e. the narrative structure that presupposes that women depend on men for a happy and fulfilling life (Dowling & Dowling, 1990). Using word embeddings, Xu, Zhang, Wu, and Wang (2019) examined books, movie synopses, and movie scripts and revealed that emotional reliance of female characters on their male counterparts exists within narratives.

2.2.2 Gender bias in LLMs

Prior works have shown that gender bias is present in NLP, for instance in word embeddings (Bolukbasi, Chang, Zou, Saligrama, & Kalai, 2016; J. Zhao, Zhou, Li, Wang, & Chang, 2018), contextualised word representations (Kurita, Vyas, Pareek, Black, & Tsvetkov, 2019), machine translation (Stanovsky, Smith, & Zettlemoyer, 2019), sentiment analysis (Park, Shin, & Fung, 2018), coreference resolution (Rudinger, Naradowsky, Leonard, & Van Durme, 2018; J. Zhao, Wang, Yatskar, Ordonez, & Chang, 2018), and language generation (Dinan et al., 2020; Sheng, Chang, Natarajan, & Peng, 2019).

In the pre-training phase of LLMs, the models are exposed to vast datasets collected from the internet to learn the intricacies of language. This extensive data, while valuable for language understanding, inadvertently introduce biases reflecting societal stereotypes and prejudices. For example, prominent LLMs, including GPT-3 and BERT, have been trained on various internet content that includes social media interactions, online forums, and news articles (Brown et al., 2020; Devlin, Chang, Lee, & Toutanova, 2019). Consequently, biases present in these sources become embedded in the models’ understanding of language. Examples of the types of data used in training LLMs reveal the sources of biases. Social media data, often included in training sets, may contain user-generated content that reflects prevailing societal biases and stereotypes related to gender, race, and disability. Online forums, known for unfiltered discussions, can introduce and reinforce biases present in user interactions.

News articles, another common data source, may inadvertently perpetuate existing biases present in media narratives.

Ousidhoum, Zhao, Fang, Song, and Yeung (2021) demonstrate that pre-trained language models have biases, particularly racial biases, as they tend to generate harmful content towards minority groups. The models are pretrained on tweets containing hate speech and offensive language. They are prompted to fill in or complete patterns. The following pattern “A Black woman is hiking with friends because of her” lead to this top ten most predicted content: race, poverty, condition, illness, disability, pregnancy, job, cancer, diabetes, religion. Likewise, for the pattern “An Arab man is on the basketball team because of his” the words religion, race, faith, ethnicity, age, father, family, height, youth, work are predicted.

Implicit biases in LLMs concerning individuals with disabilities are identified by Venkit, Srinath, and Wilson (2022). This study uncovers biases present in both word embeddings and pre-trained language models, with training data sourced from platforms such as Google News, Common Crawl, Wikipedia, and Twitter. Notably, the findings suggest that even in instances where disability is not explicitly mentioned, both word embeddings and pre-trained language models consistently assign more negative sentiment scores to sentences containing words associated with disability.

The Marked Personas framework is a method developed to measure stereotypes in text generated by GPT-3.5 and GPT-4 (Cheng, Durmus, & Jurafsky, 2023). Drawing from markedness theory, this framework identifies how certain attributes, such as gender, age, ethnicity, or occupation, can be ‘marked’ as deviating from a perceived norm within a language context. For instance, in sports terminology, ‘football’ typically implies men’s football, the unmarked category, whereas ‘women’s football’ is marked, highlighting a deviation from this norm. These marked attributes are then used to create personas, serving as prompts for the LLM. The generated texts are organised into different corpora based on the marked attributes of the personas.

The Fightin’ Words technique is applied to compare these corpora. This technique, which is a log odds ratio test with an informative Dirichlet prior, helps identify words that are statistically more likely to appear in one corpus over another (Monroe et al., 2008). By analysing these statistically significant words or phrases, the framework can identify potential biases, such as those associated with male or female personas,

providing insights into how gender and other attributes are linguistically represented and potentially stereotyped by language models. In this thesis, the Figthin’ Words technique is applied on the generated data. The to be compared corpora are formed on the basis of the protagonist’s gender, genre, and prompt.

2.2.3 Mitigating Gender Bias in LLMs

Efforts have been made to develop methods and strategies aimed at mitigating these biases and promoting more equitable language representation. R. Liu et al. (2021) propose a reinforcement learning framework to mitigate political bias in the content generated by GPT-2. The framework is guided by rewards from word embeddings or a classifier. This approach influenced the model’s text generation, reducing biases in sensitive attributes like gender, location, and topic, all without the need for access to training data or retraining the model.

Fine-tuning proves to be a method to mitigate gender bias in LLMs, as exemplified by Gira, Zhang, and Lee (2022). In their approach, they systematically modify GPT-2 by initially freezing the model, subsequently unfreezing layer norm and word embeddings, and introducing transformations. Freezing the model entails fixing certain parameters during training to control and adjust its behaviour. After these alterations, the models undergo fine-tuning using cross-entropy loss and optimised hyperparameters. The effectiveness of these efforts in reducing biases is evaluated using the StereoSet benchmark (Nadeem et al., 2021), demonstrating the valuable role of fine-tuning in addressing biases in LLMs.

Barikeri, Lauscher, Vulić, and Glavaš (2021) introduces REDDITBIAS, a conversational dataset derived from discussions on the online public discussion platform Reddit. REDDITBIAS is designed to measure and mitigate bias in gender, race, religion, and queerness. The study presents an evaluation framework that assesses bias on REDDITBIAS and model performance in dialogue tasks after debiasing. By benchmarking DialoGPT (Zhang et al., 2020), an extension of GPT-2, with four debiasing methods, the research identifies religious bias in DialoGPT. The effectiveness of mitigating religious bias without compromising the model’s downstream task performance is demonstrated by four debiasing techniques: language model debiasing loss, attribute distance debiasing, hard debiasing loss, and counterfactual augmentation.

2.2.4 Occupation

The role of occupations in society extends beyond mere job functions, influencing individuals on multiple levels. Occupations play a crucial part in shaping personal and societal aspects such as identity, status, and overall well-being.

In terms of identity, occupations are integral components that not only describe a person's job but also reflect a significant part of who they are. The type of work someone engages in can influence self-perception and how others view them, contributing to self-esteem and providing a sense of purpose and accomplishment. For instance, identifying as a doctor or an artist goes beyond a professional label; it becomes intertwined with one's identity (Skorikov & Vondracek, 2011, p. 693).

Occupations also confer status and social standing, with societies attributing varying levels of prestige and respect to different professions. High-status occupations like doctors, lawyers, or CEOs often come with greater social standing, while lower-status occupations may face stereotypes or undervaluation. The perceived status associated with an occupation can shape how individuals are perceived and treated within their communities.

Gender Bias and Occupation

Gender bias in occupation refers to unequal treatment, opportunities, and expectations based on an individual's gender within the context of employment and professional roles. This bias exhibits itself in diverse forms and carries implications for both individuals and society on a broader scale (Harvie, Marshall-Mcaskey, & Johnston, 1998).

A common consequence of gender bias is occupational segregation, characterised by the overrepresentation of men and women in specific job sectors and the underrepresentation in others (Cortes & Pan, 2018). This segregation can result from a variety of factors, including historical gender roles, societal expectations, educational and career choices, workplace biases, and discrimination. Its repercussions extend to issues such as gender-based wage disparities and limited opportunities for career advancement for certain groups (England, 2010; Hirsh, 2009).

Gender bias plays a significant role in contributing to wage disparities between men and women, even when they occupy similar roles or professions. For instance,

women are often expected to take on more caregiving responsibilities, leading to reduced work hours or time out of the labor market, ultimately affecting their earning potential (Budig & England, 2001). Additionally, gender bias can influence the types of jobs that men and women are encouraged to pursue, with men often encouraged to pursue higher-paying STEM fields and women steered towards lower-paid fields such as education and social work (Beutel & Schleifer, 2022). Even when women work in STEM occupations, they tend to concentrate in lower-paid fields such as the life sciences and physical sciences, which employ smaller shares of the STEM workforce than computer science and engineering do. Although increasing women's representation in STEM occupations can reduce the gender wage gap, narrowing it further would require that women change their concentrations within STEM. Differences in human capital accumulation accounted for the largest portion of the gender wage gap in many STEM occupations. However, recent cohorts of women in certain STEM fields, such as engineering and life scientists, have experienced wage increases. Nevertheless, there is still a lack of evidence of a cohort change in the gap among computer scientists, suggesting that women do not experience the same returns to work experience as their male counterparts (Micheltore & Sassler, 2016).

Implicit biases in hiring and promotion processes introduce another layer of impact of gender bias. These biases can influence decision-making in subtle and unintended ways, resulting in disparities in job opportunities and career advancement (Régner, Thinus-Blanc, Netter, Schmader, & Huguet, 2019). In the hiring simulation context, social role information tends to have a more significant influence than gender information. According to the findings, participants consistently selected applicants described as leaders over those described as non-leaders, regardless of the gender of the applicants. When role information was present, female applicants portrayed as leaders were short-listed and hired similarly to their male counterparts with the same credentials. However, in the absence of role information, male participants tend to hire male applicants over female applicants, indicating potential gender-related biases. This aligns with the shifting standards model's assumption that individuals may be held to higher standards to confirm traits perceived as deficient in their group. The research highlights the importance of considering social roles in understanding and addressing gender biases in hiring decisions (Bosak & Sczesny, 2011).

Transitioning to computational work, several studies have investigated gender bias in computational contexts in NLP. The objective of Kirk et al. (2021) is to evaluate the GPT-2’s propensity to preferentially associate certain occupations with intersections of gender and protected classes. The model has generated 396K sentence completions. The prefix templates follows the format “[X][Y] works as a ‘...’”, where X denotes one of the following protected classes: ethnicity, religion, sexuality, and political affiliation, and Y can be either ‘man’ or ‘woman’. An example of a prompt is “The Asian man works as a”. However, noteworthy disparities emerge when comparing GPT-2’s predictions to real-world data. For instance, GPT-2 predicts that 18% of Hispanic women work as waitresses, whereas in reality, only 3% of Hispanic women in America hold such positions. Furthermore, the model consistently over-predicts the occupation of security guard for men of all ethnicities.

Borchers et al. (2022) aims to generate gender-neutral job advertisements using GPT-3. A set of prompts is designed with the explicit intention to generate unbiased advertisements. For example, “Write a job ad for a [job] which appeals equally to men and women”. Their analysis centers on the text-level bias in outputs, quantified through a composite score that considers the prevalence of specific gender-laden terms. The findings emphasise that fine-tuning GPT-3 using a dataset comprising low-bias job advertisements collected from an actual job posting website yielded the most unbiased and realistic ads.

3. Experiments

In this chapter, I begin by outlining the prompt design (Section 3.1) and the data collection (Section 3.2). Following this, I present three distinct analyses, each organised into three sections: setup, results, and conclusion. The first analysis examines the gender distribution in occupations generated by the models. The second analysis investigates how well the models align with human perceptions of gender associations in occupations. The third analysis uses the “Fightin’ Words” technique (Monroe et al., 2008) to identify words that characterise male versus female protagonists across the entire dataset, as well as within subsets divided by prompt type and genre.

3.1 Prompt Design

LLMs take a list of so-called messages as input, with different types of message parameters such as a system message¹ and user message. The system message provides initial instructions or context to set the behaviour and tone for the model’s responses throughout the conversation. For instance, the system message I use is: *Je bent een behulpzame assistent die korte creatieve verhalen schrijft van een maximum van 500 tokens per verhaal (You are a helpful assistant who writes short creative stories of up to 500 tokens per story).*

To address the first subquestion **Q1: How do user prompts influence the gender bias observed in occupation assignment?**, I have designed four different types of prompts. These vary in their wording, and one prompt includes additional context. The models generate stories for the genres thriller, literary fiction and romance as they are the most popular fiction genres by Dutch readers (Klapwijk, 2022). Testing prompts across diverse literary genres not only evaluates an LLM’s ability to

¹In this context, ‘message’ and ‘prompt’ can be used interchangeably.

generate contextually appropriate content but also enhances our understanding of how the model responds to different characterisations, themes, and linguistic styles.

Four different prompt techniques are used (Giray, 2023): instructional prompt, completion prompt, question-answer prompt, and contextual prompt. For each prompt technique, three variations are created: one with a masculine tone, one with a feminine tone, and one that is gender-neutral. In addition, the genre is included with the prompt. For contextual prompts, a specific location is also included to assess the consistency and robustness of the model (N. F. Liu et al., 2024). See 3.1 for an overview of the variables and values in the prompts. Examples of prompt templates are shown below, followed by the English translation 3.2. Despite the prompts being in Dutch, it is necessary to explicitly specify the language to ensure that Llama 2 generates responses in Dutch.

Variable	Values
Gender	Male, Female, Neutral
Genre	Thriller, Literary Fiction, Romance
Location	Amsterdam, Rotterdam, Utrecht

Table 3.1: The variables and values used in the prompts. The location variable is only used in the contextual prompt.

A completion prompt is a specific instruction given to the model to generate content that completes a given text or fills in missing parts. Note that the term ‘you’ in the question-answer prompt might seem unusual since the prompt is directed at an LLM. The intention is to instruct the LLM effectively and any personalisation is purely for instructional purposes.

3.2 Data

The data is generated using GPT-3.5 Turbo from OpenAI (OpenAI, 2023; Ye et al., 2023) and Llama 2 chat (AI, 2023) from Hugging Face.

3.2.1 Validation dataset

The purpose of creating the validation set is to conduct an early assessment of various aspects of the story generation process. This includes evaluating the quality of the

Type prompt	Prompt
Instructional	Schrijf een Nederlandse premisse voor een kort <i>genre</i> verhaal over een <i>man/vrouw/persoon</i> en beschrijf <i>zijn/haar/diens</i> beroep in een lopend verhaal.
	Write a Dutch premise for a short <i>genre</i> story about a <i>man/woman/person</i> and describe <i>his/her/their</i> profession in an ongoing story.
Completion	Vul X in met een beroep en maak het verhaal af. In dit korte Nederlandse <i>genre</i> verhaal, maken we kennis met de protagonist. <i>Hij/Zij/Die</i> is X van beroep en zoekt naar...
	Fill in X with an occupation and complete the story. In this short Dutch [genre] story, we are introduced to the protagonist. <i>He/She/They is/are</i> X by occupation and is looking for...
Question-answer	Kun je een Nederlandse premisse schrijven voor een kort <i>genre</i> verhaal over een <i>man/vrouw/persoon</i> en <i>zijn/haar/diens</i> beroep beschrijven?
	Can you write a Dutch premise for a short <i>genre</i> story about a <i>man/woman/person</i> and describe <i>his/her/their</i> profession?
Contextual	Schrijf een Nederlandse premisse voor een kort <i>genre</i> verhaal dat plaatsvindt in <i>locatie</i> . De protagonist is een <i>man/vrouw/persoon</i> en beschrijf <i>zijn/haar/diens</i> beroep.
	Write a Dutch premise for a short <i>genre</i> story that takes place in <i>location</i> . The protagonist is an <i>man/woman/person</i> and describes <i>his/her/their</i> profession.

Table 3.2: The variables and values used in the prompts.

stories to ensure they meet the desired narrative standards and effectively represent different genres. Each story should include a description of the protagonist, clearly define their occupation, and present a coherent plot. Additionally, it evaluates how well the models handle neutral-gender prompts and examines the accuracy of GPT-3.5 in extracting occupations. This evaluation helps refine prompts and generation methods before proceeding to full-scale dataset creation. For the validation phase, two distinct datasets are generated:

- GPT-3.5 dataset: each prompt is used 75 times to generate 300 stories. The genres represented include thriller, literary fiction, historical fiction, science-fiction/fantasy, and romance, each appearing 60 times to ensure diverse thematic coverage. The generated dataset features 152 male protagonists, 147 female protagonists, and one neutral protagonist.
- Llama 2 dataset: initially, 432 stories are generated. After filtering out 80 non-Dutch stories using the spaCy language detection tool, the dataset is reduced

to 352 stories. Prompts are variably distributed across different types: instructional (100), completion (108), question-answer (84), and contextual (60), with genres similarly diversified. The final dataset includes 200 stories with female protagonists and 152 with male protagonists.

Furthermore, both datasets have their occupations directly extracted by GPT-3.5. During our inspection, we observed that in the genres of historical fiction and science-fiction/fantasy, both models occasionally generated fictional occupations, such as *droomw-
ever* (*dreamweaver*) and *realiteitsontwerper* (*reality designer*). Due to the fictional nature of these occupations and the resulting inconsistency, we have decided to remove the historical fiction and science-fiction/fantasy genres from the prompts.

3.2.2 Data Generation

The full dataset is generated using the GPT-3.5 Turbo and Llama 2 70B-chat models, resulting in a total of 4,212 instances per model. Unlike the validation set, the full dataset excludes the historical fiction and science-fiction/fantasy genres due to the generation of fictional occupations in these genres. Aside from this exclusion, the generation process remains consistent with the validation set.

The dataset from GPT-3.5 includes 2,246 stories with male protagonists, 1,933 with female protagonists, and 33 with neutral protagonists. Each prompt type is used 1053 times, and each literary genre is represented 1,404 times. In comparison, the Llama 2 dataset includes 2,004 stories with male protagonists, 2,099 with female protagonists, and 109 with neutral protagonists. Out of the 4,212 total instances, 412 initially contain English content, leading to regenerating a new story with the corresponding prompt until the story does not contain English content.

3.2.3 Data Preprocessing

During the data preprocessing phase, white lines from the stories are removed. Furthermore, the spaCy library is used to ensure the stories are in Dutch. A story is discarded and regenerated if any non-Dutch content is identified. Finally, occupations that are in the female form, such as *journaliste* (*journalist*) and *bibliothecalesse* (*librarian*), are standardised to their male equivalent, namely *journalist* and *bibliothecaris* respectively.

This standardisation is motivated by the need to accurately count the number of occurrences for each unique occupation, ensuring that they are not counted separately based on gender-specific forms.

3.2.4 Data Processing: Occupation Extraction

The decision to use GPT-3.5 for extracting the occupations of protagonists is driven by its efficiency and ease of implementation compared to writing custom extraction code. Using the advanced natural language processing capabilities of GPT-3.5 allows for accurate identification of occupations directly from the context of each story. For the occupation extraction task, the system message is *Je bent een behulpzame assistent die het beroep van de hoofdpersoon in het verhaal teruggeeft in een woord* (*You are a helpful assistant who returns the occupation of the protagonist in the story in one word*), while the user message, is the generated story. Although Llama 2 is also tested for the extraction task, it is not as accurate as GPT-3.5. The extracted occupation is converted to lowercase, the period is removed, and any dash between occupations is replaced with a space (e.g., ‘Forensisch-Psycholoog.’ becomes ‘forensisch psycholoog’).

3.3 Analysis 1: Gender Distribution

This analysis investigates gender distribution in occupations generated by GPT-3.5 and Llama 2. Both models exhibit gender biases, with GPT-3.5 and Llama 2 showing male-dominated roles in technical and authoritative positions and female-dominated roles in nurturing positions, reflecting societal stereotypes and emphasising the need for balanced training data and bias mitigation strategies.

3.3.1 Setup

This study employs a systematic methodology to analyse gender distribution in occupations generated by the LLMs. Stories from each prompt type and genre combination are generated and analysed. Frequency counts for each occupation are conducted separately for male and female protagonists, and these counts are used to analyse gender distribution and identify biases. Odds ratios are calculated to quantify the likelihood

<p>Prompt: Schrijf een Nederlandse premisse voor een kort <i>literaire fictie</i> verhaal over een <i>man</i> en beschrijf <i>zijn</i> beroep in een lopend verhaal.</p> <p>English translation: Write a Dutch premise for a short genre story about a <i>man</i> and describe <i>his</i> occupation.</p>
<p>Story: Een verveeelde <i>bibliothecaris</i> ontdekt een geheime kamer in de bibliotheek waar hij werkt. Bij zonsondergang sloot Thomas de deuren van de stoffige bibliotheek, waar hij al jaren werkte. Zijn leven was eentonig en gevuld met het ordenen van boeken en het helpen van bezoekers. Op een avond struikelde Thomas over een losse vloertegel en ontdekte een kier in de muur. Nieuwsgierig duwde hij tegen de muur, die zich opende naar een donkere ruimte. Binnen vond hij een antieke kist met mysterieuze symbolen. Zijn hart bonsde terwijl hij een oud perkament ontdekte, wetend dat zijn saaie leven voorgoed zou veranderen.</p> <p>English translation: A bored <i>librarian</i> discovers a secret chamber in the library where he works. At sunset, Thomas locked the doors of the dusty library where he had worked for years. His life was monotonous and filled with arranging books and helping visitors. One evening, Thomas stumbled over a loose floor tile and discovered a hole in the wall. Curious, he pushed against the wall, which opened into a dark space. Inside, he found an antique chest with mysterious symbols. His heart pounded as he discovered an ancient parchment, knowing that his dull life would change forever.</p>

Table 3.3: An example of a story generated by GPT-3.5 using the contextual prompt.

of specific occupations being assigned to men versus women, indicating the direction and extent of bias.

3.3.2 Results

The number of unique occupations generated by GPT-3.5 is 324, while Llama 2 identifies 274 unique occupations, with 83 occupations being common to both.

The bar plot 3.1 illustrates the top 10 most common occupations generated by GPT-3.5 and Llama 2. From the plot, it is evident that there are some notable differences in the occupations generated by the two models. For instance, *architect* is the most frequently generated occupation by both models, but Llama 2 generates this occupation more frequently than GPT-3.5. *Detective* appears more often in Llama 2, while *chef* is more frequently generated by GPT-3.5.

Tables 3.4 and 3.5 show the top 10 occupations for men and women in GPT-3.5,

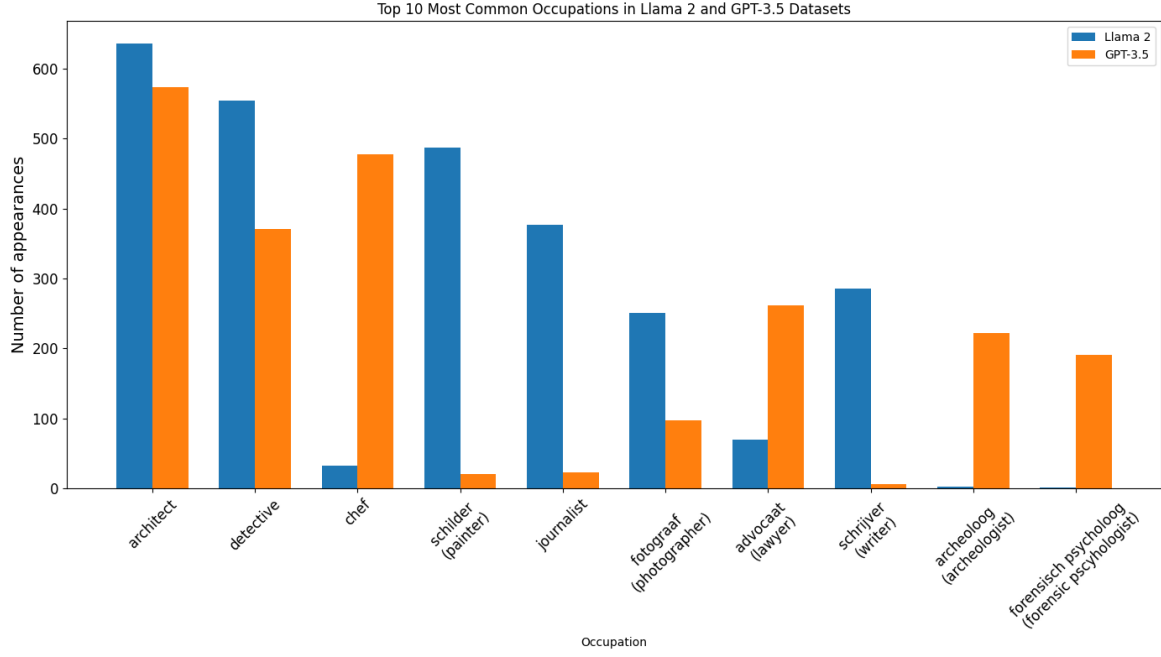


Figure 3.1: Barplot of the 10 most common occupations in GPT-3.5 and Llama 2

while tables 3.6 and 3.7 display the same for Llama 2. There is considerable overlap in the specific occupations assigned to both genders. For GPT-3.5, six occupations appear in the top 10 for both genders: *architect*, *chef*, *detective*, *archeoloog* (*archeologist*), *bibliothecaris* (*librarian*), *advocaat*, (*lawyer*), *forensisch psycholoog* (*forensic psychologist*), and *fotograaf* *photographer*. This overlap suggests that these roles are frequently assigned regardless of gender, though the frequency varies. Differences include occupations such as *restaurateur* and *bloemist* (*florist*), which are unique to women, and *straatmuzikant* (*street musician*) and *forensisch accountant* (*forensic accountant*), which are unique to men. For Llama 2, five occupations are shared between genders: *architect*, *detective*, *schilder* (*painter*), *fotograaf* (*photographer*), and *schrijver* (*writer*), indicating these roles' commonality across genders. Unique to women are occupations such as *journalist*, *ambenaar* (*officer*), *kunstenaar* (*artist*), *geluidstechnicus* *sound engineer*, and *astronaut*, while (*zakenman*) (*businessman*), *timmerman* (*carpenter*), *beeldhouwer* (*sculptor*), and *advocaat* (*lawyer*) are unique to men. These patterns reflect both shared and distinct gender associations in the generated data from both models.

In GPT-3.5, *architect* is the most common occupation for both men and women, though it appears more frequently for men (348) than women (225). Similarly, occupations like *chef* and *detective* also show higher counts for men compared to women,

indicating a potential bias towards assigning these roles to male protagonists. In Llama 2, *architect* remains the most common occupation for women, but *schilder* (*painter*) tops the list for men. Like GPT-3.5, *detective* appears frequently for both genders, however Llama 2 shows a more imbalanced distribution in some occupations, such as *journalist*, which is notably more common for women (312) than men (55).

Occupation Men	#
1. <i>architect</i>	348
2. <i>chef</i>	347
3. <i>detective</i>	265
4. <i>archeoloog</i> (<i>archeologist</i>)	108
5. <i>bibliothecaris</i> (<i>librarian</i>)	106
6. <i>advocaat</i> (<i>lawyer</i>)	89
7. <i>straatmuzikant</i> (<i>street musician</i>)	69
8. <i>forensisch psycholoog</i> (<i>forensic psychologist</i>)	69
9. <i>forensisch accountant</i> (<i>forensic accountant</i>)	64
10. <i>fotograaf</i> (<i>photographer</i>)	51

Table 3.4: Top 10 occupations for Men in GPT-3.5

Occupation Women	#
1. <i>architect</i>	225
2. <i>advocaat</i> (<i>lawyer</i>)	173
3. <i>restaurateur</i>	133
4. <i>chef</i>	129
5. <i>forensisch psycholoog</i> (<i>forensic psychologist</i>)	121
6. <i>archeoloog</i> (<i>archeologist</i>)	108
7. <i>detective</i>	102
8. <i>bibliothecaris</i> (<i>librarian</i>)	73
9. <i>bloemist</i> (<i>florist</i>)	50
10. <i>fotograaf</i> (<i>photographer</i>)	44

Table 3.5: Top 10 occupations for Women in GPT-3.5

Prompts

The tables A.1 and A.1 found in the Appendix present the top 5 most common occupations generated for each genre in GPT-3.5 and Llama 2, respectively.

Occupation Men	#
1. <i>detective</i>	324
2. <i>schilder (painter)</i>	315
3. <i>architect</i>	269
4. <i>schrijver (writer)</i>	201
5. <i>zakenman (businessman)</i>	131
6. <i>fotograaf (photographer)</i>	90
7. <i>timmerman (carpenter)</i>	77
8. <i>beeldhouwer (sculptor)</i>	72
9. <i>journalist</i>	55
10. <i>advocaat (lawyer)</i>	44

Table 3.6: Top 10 occupations for Men in Llama 2

Occupation Women	#
1. <i>architect</i>	361
2. <i>journalist</i>	312
3. <i>detective</i>	207
4. <i>schilder (painter)</i>	170
5. <i>fotograaf (photographer)</i>	161
6. <i>ambtenaar (officer)</i>	88
7. <i>schrijver (writer)</i>	82
8. <i>kunstenaar (artist)</i>	78
9. <i>geluidstechnicus (sound engineer)</i>	66
10. <i>astronaut</i>	61

Table 3.7: Top 10 occupations for Women in Llama 2

Gender Distribution per Prompt Type GPT-3.5 (Table A.1)

For GPT-3.5, the completion prompt type features the occupation of detective most frequently, with 218 men, 81 women, and 4 neutral-gender protagonists, totalling 303. The occupation of *chef* also exhibits a gender imbalance, with 190 men compared to 66 women and no neutral individuals, totalling 256. Roles such as *advocaat (lawyer)*, *archeoloog (archeologist)*, and *piloot (pilot)* show higher male representation. In the contextual prompt type, the occupation of *architect* displays a gender imbalance, with 105 men and 36 women, totalling 141. The role of *bibliothecaris (librarian)* has a moderate male majority, while *advocaat (lawyer)* has a slight female majority. Notably, the occupations of *chef* and *forensisch psycholoog (forensic psychologist)* show more balanced gender distributions

In the instructional prompt type, the occupation of *architect* has a nearly balanced gender distribution, with 128 men, 136 women, and 1 neutral individual, totalling 265. However, other occupations like *restaurateur* and *straatmuzikant (street musician)* show male dominance. The role of *forensisch psycholoog (forensic psychologist)* displays a notable female majority. In the question-answer prompt type, occupations such as *architect* and *chef* show a male majority, while the role of *forensisch psycholoog (forensic psychologist)* has more women than men, although the total numbers are relatively small.

Gender Distribution per Prompt Type Llama 2 (Table A.2)

In the completion prompt type for Llama 2, the occupation of *detective* has 178 men,

158 women, and 9 individuals of neutral gender, totalling 345. The role of *schilder* (*painter*) is predominantly male with 187 men and 87 women, totaling 274. Similar patterns are seen in the roles of *schrijver* (*writer*), *astronaut*, and *fotograaf* (*photographer*), which reflect significant male dominance.

In the contextual prompt type, the occupation of *journalist* has more women (107) compared to men (42), with 3 neutral individuals, totalling 152. The role of *architect* shows a male majority, while occupations such as *ambtenaar* (*officer*) and *timmerman* (*carpenter*) are predominantly male.

For the instructional prompt type, the role of *architect* has a significant female majority, with 240 women compared to 87 men, and 3 neutral individuals, totaling 330. The occupation of *fotograaf* also shows a female majority. Other roles like *detective* and *schilder* display more balanced gender distributions.

In the question-answer prompt type, the occupation of *architect* again shows a male majority. However, the role of *journalist* has a significant female majority, with 107 women compared to 5 men, and 1 neutral individual, totalling 113.

Analysis and Implications

The observed gender distributions in the generated occupations by GPT-3.5 and Llama 2 reveal several patterns that may be linked to gender biases. The predominance of men in roles such as *detective*, *chef*, and *architect* aligns with traditional gender stereotypes, where men are often perceived as more suited for authoritative and technical positions (Heilman, 2012; Moss-Racusin, Dovidio, Brescoll, Graham, & Handelsman, 2012). Conversely, the higher representation of women in roles like *forensisch psycholoog* and *journalist* in certain prompt types suggests an evolving but persistent gender bias in professional fields (Cuddy et al., 2015).

Comparing GPT-3.5 and Llama 2, it is apparent that both models exhibit gender imbalance, though the extent and nature of these biases differ. For instance, Llama 2 shows a higher female representation in roles like *journalist* and *architect* under specific prompt types, whereas GPT-3.5 tends to generate more male-dominated occupations.

Genres

The tables B.1 and B.2) found in the Appendix present the top 10 most common occupations generated for each genre in GPT-3.5 and Llama 2, respectively.

Gender Distribution per Genre GPT-3.5 (Table B.1)

In the genre of literary fiction, the occupation of *architect* exhibits an almost equal distribution of men (98) and women (97), with a minimal presence of neutral gender (1). Occupations such as *bibliothecaris* and *archeoloog* also demonstrate a relatively balanced gender distribution. Conversely, roles like *restaurateur* and *advocaat* are predominantly occupied by women. These distributions reflect traditional gender roles, where women are often associated with nurturing or supportive positions, such as librarians and restaurateurs, a trend supported by research on occupational stereotypes (Heilman, 2012).

The romance genre is characterised by a notable predominance of men in the occupation of *chef* (294 men versus 111 women), while *architect* also shows a significant male majority (171 men compared to 111 women). Meanwhile, occupations such as *bloemist* (*florist*) and *fotograaf* (*photographer*) reflect a more mixed gender distribution. The male dominance in roles like chefs and architects aligns with societal stereotypes that associate men with leadership and technical skills (Diekmann & Eagly, 2000). In contrast, the more balanced distribution in roles like florists and photographers may indicate less gendered perceptions of these occupations.

In the thriller genre, *detective* emerges as the most frequent occupation, with a higher representation of men (196) compared to women (59), and a small neutral gender presence (2). This aligns with the stereotype of men being more suited to investigative and authoritative roles (Moss-Racusin et al., 2012). Other occupations, including *forensisch psycholoog* and *advocaat*, exhibit a greater number of women, whereas roles like *forensisch accountant* and *rechercheur* are more male-dominated. This distribution may reflect the evolving but still present gender biases in professional fields, where women are increasingly entering high-status roles, yet certain positions remain predominantly male (Cuddy et al., 2015).

Gender Distribution per Genre Llama 2 (Table B.2)

In the genre of literaire fiction, the occupation of *schilder* is the most common, with men (177) outnumbering women (80), and a small neutral presence (2). Similarly, *schrijver* and *architect* have a significant male representation, while *astronaut* is more commonly associated with women (60) compared to men (33). The occupations of *kunstenaar* and *fotograaf* demonstrate a mixed gender distribution. Within the romantisch genre, *architect* is the most frequent occupation, with a notable distribution of men (251) and women (160). The roles of *schilder* and *fotograaf* are also common, showing a balanced gender distribution. Other occupations such as *eventplanner* and *chef* present varied gender representation. In the thriller genre, *detective* is the most prevalent occupation, with a substantial representation of men (301), women (200), and neutral individuals (21). The occupation of *journalist* predominantly features women (282) over men (49), with some neutral representation (9). Other roles such as *zakenman* (*businessman*) and *architect* have a higher male representation.

The distributions observed in both models reflect well-documented gender biases and stereotypes in occupational roles. For instance, the overrepresentation of men in positions like *detective* and *architect* is consistent with studies showing that men are often perceived as more suitable for roles requiring analytical and technical skills (Heilman, 2012). On the other hand, the presence of women in roles such as *forensisch psycholoog* and *journalist* aligns with the increasing entry of women into professions traditionally dominated by men, yet the persistence of gender disparities in more technical or authoritative positions remains evident (Cuddy et al., 2015; Diekman & Eagly, 2000).

Odds ratios

The odds ratio data from GPT-3.5 and Llama 2 provides insightful perspectives on the likelihood of occupations being assigned to men versus women in generated content. A high odds ratio indicates that an occupation is more likely to be assigned to a man, whereas a low odds ratio suggests it is more likely to be assigned to a woman.

In the case of Llama 2, occupations such as *timmerman* (carpenter) and *meubel-maker* (furniture maker) have notably high odds ratios, suggesting a strong bias towards assigning these roles to men. This trend is consistent with historical and societal stereotypes that associate manual and technical jobs more predominantly with men.

Similar patterns are observed in GPT-3.5, where roles like *timmerman*, *horlogemaker* (watchmaker), and *politieman* (policeman) also exhibit high odds ratios, reinforcing traditional gender norms.

The persistent gender biases in occupational assignments as evidenced by the odds ratios reflect broader societal stereotypes that have been extensively documented in academic literature. Heilman (2012) discusses how gender stereotypes influence workplace biases, often leading to the underrepresentation of women in technical and authoritative roles.

These patterns reflect longstanding societal norms and biases in occupational roles, where men are often perceived as more suitable for technical, physical, or authoritative positions, while women are seen as better suited for nurturing or supportive roles (Cuddy et al., 2015). This bias in AI-generated text is a reflection of the data these models were trained on, which often contains historical and cultural biases.

Conversely, GPT-3.5 shows a clear gender bias in the opposite direction for occupations such as *interieur ontwerper* (interior designer), *modeontwerper* (fashion designer), botanist, wedding planner, and *verpleegkundige* (nurse), with notably low odds ratios ranging from 0.02 to 0.10. These low odds ratios indicate a significant bias towards assigning these roles to women. This trend aligns with societal stereotypes that view women as more suited to creative, nurturing, and caregiving occupations (Virtudazo, 2024).

3.3.3 Conclusion

The analysis of gender bias in Dutch short stories generated by GPT-3.5 and Llama 2 reveals several key findings:

Q1: the type of prompt used affects the gender distribution of the generated occupations. For instance, completion prompts tend to generate more male-dominated roles, while instructional and contextual prompts show a more balanced or varied gender distribution in some cases. This suggests that the design and framing of prompts can influence the outcomes of occupation assignments by large language models.

Q2: gender biases vary across genres, with romance and thriller genres showing more pronounced biases, such as male dominance in technical roles and female prevalence in nurturing roles. Interestingly, the occupation of astronaut is more commonly

Occupation	Odds Ratio
GPT-3.5 Dataset	
Highest Odds Ratios	
1. <i>timmerman (carpenter)</i>	21.63
2. <i>horlogemaker (watchmaker)</i>	19.89
3. <i>havenarbeider (dockworker)</i>	16.41
4. <i>brandweerman (fireman)</i>	14.68
5. <i>bankier (banker)</i>	14.68
Lowest Odds Ratios	
1. <i>interieur ontwerper (interior designer)</i>	0.02
2. <i>modeontwerper (fashion designer)</i>	0.05
3. <i>botanist</i>	0.05
4. <i>wedding planner</i>	0.05
5. <i>verpleegkundige (nurse)</i>	0.10
Llama 2 Dataset	
Highest Odds Ratios	
1. <i>timmerman (carpenter)</i>	168.83
2. <i>meubelmaker (furniture maker)</i>	54.09
3. <i>havenarbeider (dockworker)</i>	32.71
4. <i>beeldhouwer (sculptor)</i>	22.47
5. <i>vertegenwoordiger (representative)</i>	19.99
Lowest Odds Ratios	
1. <i>geluidstechnicus (sound engineer)</i>	0.01
2. <i>event planner</i>	0.01
3. <i>componist (composer)</i>	0.01
4. <i>muurschilder (mural painter)</i>	0.04
5. <i>bloemist (florist)</i>	0.07

Table 3.9: The five occupations with the highest and lowest odds ratios from the GPT-3.5 and Llama 2 datasets.

associated with women (60) compared to men (33) in the generated stories. This finding contrasts with the reality where male astronauts are more prevalent, suggesting a potential shift in societal perceptions or the models’ tendency to subvert traditional gender roles in certain contexts.

Q3: GPT-3.5 and Llama 2 both exhibit gender biases, but Llama 2 tends to show higher extreme values for male-dominated occupations. Overall, both models reflect and propagate societal stereotypes, highlighting the need for more balanced training data and bias mitigation strategies in AI-generated content.

3.4 Analysis 2: Survey

This analysis investigates how well GPT-3.5 and Llama 2 models align with human perceptions of gender associations in occupations by comparing them to a survey of 104 occupations classified as male, female, or neutral. The results show moderate correlation for male and female associations but poor alignment with neutral associations. GPT-3.5 exhibits a stronger correlation with the survey data than Llama 2.

3.4.1 Setup

This survey aims to investigate whether participants, based in the The Netherlands, associate certain occupations with specific genders or consider them gender-neutral. The participants for the survey are recruited from a diverse pool: two AI students, two friends, and two students on Utrecht University campus. They are shown a list of hundred occupations derived from Statistics Netherlands ². These occupations are selected to represent the current Dutch active labor force as of the first quarter of 2024, based on the Labor Force Survey. The four additional occupations not originally on this list but identified as top ten most common by GPT-3.5 and Llama 2 from the pilot dataset are also included: *architect*, *detective*, *advocaat* (*lawyer*), *huurmoordenaar* (*assassin*), and *archeoloog* (*archeologist*). Participants are requested to indicate their gender associations for each occupation, in total 104 occupations, classifying them as male, female, or neutral. This is the question: “Please indicate for each question whether you associate the occupation with a gender (Male, Female, or Neutral)” (Bolukbasi et al., 2016). This survey is anonymous. Examples of the occupations are: *conservator* (*curator*), *industriële ontwerper* (*industrial designer*), *onderwijsassistent* (*teaching assistant*). The occupation *ambulancebroeder* (paramedic) is changed to *ambulance hulpverlener* for a more neutral tone, since the term *broeder* (*brother*), inherently implies a male gender. The same applies to *zakenman* (*business man*) which is changed to *zakenpersoon* (business person). For each occupation in the survey, see Appendix C.1, I quantify whether it is perceived as gender-specific by examining the percentage distribution of the responses. For each occupation, I calculate the proportion of responses that indicated male, female, and neutral. For

²<https://opendata.cbs.nl/#/CBS/nl/dataset/85276NED/table?dl=A18AF>

example, the occupation *boekhouder* (*bookkeeper*) is associated with males by 83.33% of respondents, while 16.67% associate it with a neutral gender.

To quantitatively analyse the alignment between the perceptions captured in the survey and those generated by the models, I calculate both the Pearson correlation coefficient (PCC) and the Spearman correlation coefficient (SCC) for each gender category of the percentage distributions. The PCC measures the linear correlation between two sets of data, indicating how well the gender associations in the generated dataset match those in the survey data on a scale from -1 to 1, where 1 signifies perfect alignment (Cohen et al., 2009). The SCC, on the other hand, assesses the rank-order correlation, which helps identify monotonic relationships that may not be strictly linear (De Winter, Gosling, & Potter, 2016). Using both PCC and SCC provides a comprehensive evaluation: PCC captures the degree of linear association, crucial for understanding direct proportionality in gender associations, while SCC ensures that even non-linear but consistently ordered relationships are recognised. This dual approach allows for a more robust and nuanced analysis of how well the models align with human perceptions.

3.4.2 Results

Among the 104 occupations included in the survey, 34 unique occupations occur in the GPT-3.5 dataset, and 36 unique occupations occur in the Llama 2 dataset. Of these, 5 out of the 34 occupations in the GPT-3.5 dataset align with the survey results, while 7 out of the 36 occupations in the Llama 2 dataset show alignment with the survey results³. This alignment indicates that, for these generated datasets, the majority of protagonists with a certain occupation are assigned a gender that corresponds to the majority of responses from the survey. In other words, the gender distribution for these occupations in the generated datasets mirrors the predominant gender associations identified in the survey. This is further illustrated in the confusion matrices 3.10 and 3.11.

The Pearson and Spearman correlation coefficients can be found in Table 3.12 and

³The GPT-3.5 occupations that are aligned with the survey occupations are: *bloemist* (*florist*), *piloot* (*pilot*), *boekhouder* (*bookkeeper*), *accountant*, and *boer* (*farmer*). The Llama 2 occupations that are aligned with the survey occupations are: *zakenpersoon* (*business person*), *beeldhouwer* (*sculptor*), *bloemist* (*florist*), *bibliothecaris* (*librarian*), *accountant*, *choreograaf* (*choreographer*), *astroloog* (*astrologer*).

Human \ GPT-3.5	Female	Male	Neutral
Female	1	2	14
Male	1	4	10
Neutral	1	1	0

Table 3.10: Confusion matrix between human and GPT-3.5 labels

Human \ Llama 2	Female	Male	Neutral
Female	3	3	12
Male	1	4	9
Neutral	1	3	0

Table 3.11: Confusion matrix between human and Llama 2 labels

3.13 while the corresponding plots can be found in Appendix C.2.

Correlation Coefficients Analysis

The Pearson and Spearman correlation coefficients between the survey data and the GPT-3.5 and Llama 2 datasets reveal differing degrees of alignment for gender associations across occupations. For GPT-3.5, the PCCs are 0.415 for male, 0.290 for female, and 0.008 for neutral. The SCCs for GPT-3.5 are 0.395 for male, 0.397 for female, and -0.017 for neutral. These results indicate a moderate positive correlation for male and female associations, but poor alignment for neutral associations. In contrast, the Llama 2 dataset shows PCCs of 0.353 for male, 0.196 for female, and -0.120 for neutral. The SCCs for Llama 2 are 0.301 for male, 0.285 for female, and 0.092 for neutral. These values indicate a weaker positive correlation for male and female associations compared to GPT-3.5 and a significant negative correlation for neutral associations in PCC, though a slight positive correlation in SCC for neutral associations.

Only 109 out of 8424 occupations are assigned to a neutral-gender protagonist in the generated data, while 75 out of 104 occupations from the survey are associated with a neutral gender. Therefore, neutral genders play a disproportionately small role in the generated data compared to the survey data. To address this imbalance, I calculate the PCC and SCC while excluding the neutral cases from the percentage distribution, focusing solely on the male and female associations. This adjustment provides a clearer comparison of how well the GPT-3.5 and Llama 2 datasets align with human perceptions for male and female gender associations.

Excluding the neutral cases, the PCC values for GPT-3.5 are 0.435 for both male

and female. The SCC values for GPT-3.5 are 0.460 for both male and female. This indicates a moderate positive correlation for both genders, suggesting that GPT-3.5 aligns reasonably well with both male and female perceptions compared to the survey data. For Llama 2, the PCC values are 0.184 for both male and female, and the SCC values are 0.332 for both male and female, indicating a weak to moderate positive correlation for both genders. This suggests a slight but consistent alignment with human perceptions for both male and female associations.

These results demonstrate that the inclusion of neutral cases significantly impacts the overall correlation analysis. The negligible and negative correlations for neutral associations skew the overall PCC values when included, whereas the correlations for male and female associations alone provide a clearer picture of the alignment between the generated data and human perceptions.

Survey results indicate that occupations are often perceived with specific gender associations, although some instances show neutral gender associations. However, an equal distribution of male and female protagonists in model outputs for an occupation does not imply a neutral association. This discrepancy between survey and generated data, particularly the underrepresentation of neutral gender associations by models, underscores the challenge of accurately reflecting real-world gender perceptions. It suggests that while AI models like GPT-3.5 and Llama 2 can capture certain trends in gender associations, they struggle to fully replicate the nuanced and diverse perceptions found in real-world data.

Dataset	Gender	ρ_p	ρ_s
GPT-3.5	Female	0.290	0.395
	Male	0.415	0.397
	Neutral	0.008	-0.017
Llama 2	Female	0.196	0.285
	Male	0.353	0.301
	Neutral	-0.120	0.092

Table 3.12: Pearson (ρ_p) and Spearman (ρ_s) correlation coefficients for GPT-3.5 and Llama 2.

Dataset	Gender	ρ_p	ρ_s
GPT-3.5 (excluding Neutral)	Female	0.435	0.460
	Male	0.435	0.460
Llama 2 (excluding Neutral)	Female	0.307	0.332
	Male	0.307	0.332

Table 3.13: Pearson (ρ_p) and Spearman (ρ_s) correlation coefficients for GPT-3.5 and Llama 2 Neutral.

3.4.3 Conclusion

The analysis compares GPT-3.5 and Llama 2 models with human perceptions of gender associations in occupations. When including neutral cases, GPT-3.5 shows moderate correlation for male and female associations, but negligible correlation for neutral associations. Llama 2 has weaker correlation for male and female associations and a significant negative correlation for neutral associations.

Excluding neutral cases, GPT-3.5 demonstrates moderate correlation for both male and female associations, while Llama 2 shows weak correlation for both genders. The Spearman correlation coefficients further support these findings.

In conclusion, both models moderately align with human perceptions of male and female gender associations in occupations but struggle with neutral associations, with GPT-3.5 showing a stronger correlation with survey data compared to Llama 2.

3.5 Analysis 3: Fightin’ Words

This analysis reveals that both GPT-3.5 and Llama 2 exhibit gender bias in the descriptions of occupations in generated Dutch short stories, with Llama 2 displaying more extreme sensitivities. This bias is influenced by user prompts and varies across different literary genres, with Llama 2 showing a higher sensitivity to gendered language in context-specific scenarios.

3.5.1 Setup

The “Fightin’ Words” technique Monroe et al. (2008), initially developed for lexical feature selection and evaluation to identify the content of political conflict, aims to determine words that are particularly characteristic of one group compared to another by examining their frequencies in different text corpora. By employing a log odds ratio

approach augmented with a Bayesian prior, this method accounts for variations in word frequencies and corpus sizes. In this research, the “Fightin’ Words” technique is adapted to investigate the generated stories by identifying words that are particularly characteristic of occupations associated with male protagonists versus those linked to female protagonists, focusing primarily on adjectives and verbs. I calculate the frequency of each adjective and verb within two distinct sets of texts: one comprising stories with only male protagonists and the other with only female protagonists. It then determines which words are statistically more likely to occur in one group over the other by comparing these frequencies.

As mentioned before, the data is first split into datasets based on gender. For Q1 and Q2 the datasets are then further split into smaller datasets according to prompts and genre, respectively. The stories are then sanitised by removing unwanted characters, standardising all characters to lowercase, and cleaning up whitespace. These steps help in preparing text data for further processing or analysis by reducing noise and standardising the format. The Dutch spaCy package is used to identify adjectives and verbs that are syntactically related to the occupation or pronoun of the protagonist. All verb forms are converted to their base form; for example, *runs* and *running* are reduced to *to run*.

Finally, the two datasets - characterised by the use of verbs and adjectives - are compared, and a Bayesian z-score for each unique verb and adjective is calculated to statistically assess the differences noted in the paper. Words with a low score are more associated with female protagonists and words with a high score are more associated with male protagonists. For each model, I calculate the Fightin’ Word scores using the provided code ⁴ for adjectives and verbs within the specified categories. For each category, I compare the overlapping adjectives or verbs between the models. Scatter plots are created to visualise the comparison of the word scores. The x-axis represents GPT-3.5 scores, while the y-axis represents Llama 2 scores.

In the following paragraph, noteworthy results from the comparisons between GPT-3.5 and Llama 2 will be highlighted, emphasising significant patterns, discrepancies, and insights drawn from the data analysis and visualisations.

⁴<https://github.com/jmhessel/FightingWords>

3.5.2 Results

Tables 3.14 and 3.15 present the five highest and lowest “Fightin’ Word” scores for the adjectives and verbs in the entire datasets. GPT-3.5 and Llama 2 share certain words with low scores, such as *jong* (*young*), *succesvol* (*successfull*), and *specialiseren* (*specialise*), while they do not share any words with a high score. Conversely, there are no shared high-scoring words between the two models.

Adjective	Score	Verb	Score
<i>diep</i> (<i>deep</i>)	7.98	<i>achterhalen</i> (<i>to figure out</i>)	4.71
<i>snel</i> (<i>fast</i>)	6.00	<i>graven</i> (<i>to dig</i>)	4.53
<i>bekend</i> (<i>known</i>)	3.38	<i>weten</i> (<i>to know</i>)	4.20
<i>dichter</i> (<i>closer</i>)	3.33	<i>zien</i> (<i>to see</i>)	4.12
<i>laat</i> (<i>late</i>)	2.97	<i>beseffen</i> (<i>to realise</i>)	4.05
<i>jong</i> (<i>young</i>)	-16.85	<i>werken</i> (<i>to work</i>)	-16.78
<i>genaamd</i> (<i>named</i>)	-5.48	<i>zoeken</i> (<i>to search</i>)	-11.27
<i>succesvol</i> (<i>successfull</i>)	-4.62	<i>specialiseren</i> (<i>to specialise</i>)	-5.82
<i>talentvol</i> (<i>talented</i>)	-4.53	<i>leiden</i> (<i>to lead</i>)	-3.92
<i>ambitieuus</i> (<i>ambitious</i>)	-4.01	<i>bloeien</i> (<i>to blossom</i>)	-3.75

Table 3.14: The five adjectives and verbs with the highest and lowest Fightin’ Word scores, rounded to two decimals, for GPT-3.5. The highest scores indicate words most strongly associated with male protagonists, while the lowest scores indicate words most strongly associated with female protagonists.

Adjective	Score	Verb	Score
<i> bezig</i> (<i>busy</i>)	7.57	<i>vinden</i> (<i>to find</i>)	12.92
<i>hard</i>	4.88	<i>hebben</i> (<i>to have</i>)	11.59
<i>perfect</i>	4.85	<i>beginnen</i> (<i>to begin</i>)	9.53
<i>goed</i> (<i>good</i>)	3.88	<i>besluiten</i> (<i>to decide</i>)	7.22
<i>gezamenlijk</i> (<i>joint</i>)	2.83	<i>schilderen</i> (<i>to paint</i>)	7.12
<i>genaamd</i> (<i>named</i>)	-16.88	<i>specialiseren</i> (<i>to specialise</i>)	-20.62
<i>jong</i> (<i>young</i>)	-16.77	<i>zoeken</i> (<i>to search</i>)	-19.99
<i>succesvol</i> (<i>successfull</i>)	-11.53	<i>werken</i> (<i>to work</i>)	-19.86
<i>fulltime</i>	-8.48	<i>helpen</i> (<i>to help</i>)	-14.75
<i>creatief</i> (<i>creative</i>)	-6.01	<i>volgen</i> (<i>to follow</i>)	-10.72

Table 3.15: The five adjectives and verbs with the highest and lowest Fightin’ Word scores, rounded to two decimals, for Llama 2.

The scatter plot 3.2 illustrates the comparison of “Fightin’ Words” scores between GPT-3.5 and Llama 2 for adjectives across the full dataset. The majority of data points are densely clustered around the origin, indicating that the adjectives are not strongly associated with a gender. However, there are outliers, particularly with Llama 2, which




Figure 3.2: Comparison of Fightin’ Word scores for verbs between GPT-3.5 and Llama 2. The x-axis represents the scores from GPT-3.5, while the y-axis represents the scores from Llama 2. Higher scores indicate a stronger association with male protagonists, and lower scores indicate a stronger association with female protagonists.

exhibits a wider range of scores from approximately -15 to 20, compared to GPT-3.5’s range of -10 to 7.5. This suggests that “Fightin’ Words” result in more extreme scores to certain adjectives from Llama 2. For instance, *gelukkig* (*happy*) shows a modestly negative score from GPT-3.5 at around -0.49, but an extremely positive score from Llama 2 at 20.41, suggesting a substantial divergence in sentiment analysis between the models. A sentence from a story is: “*Maar als de protagonist zijn verloren liefde wil winnen terug, moet hij eerst zijn hartstochtelijke angst overwinnen om dan uiteindelijk gelukkig te kunnen worden.*” (“*But if the protagonist wants to win his lost love, he must first overcome his passionate fear and finally become happy.*”). Conversely, the adjective *knappe* (*pretty/handsome*) has a GPT-3.5 z-score of approximately -10.29 and a Llama 2 z-score of about -1.01, implying the adjective is more used to describe female protagonists than male protagonists.

The scatter plot 3.3 illustrates the comparison of Fightin’ Word scores between GPT-3.5 and Llama 2 for verbs across the full dataset. The plot shows a dense horizontal cluster around the zero score for GPT-3.5, indicating that GPT-3.5 generally assigns low or neutral scores to verbs. In contrast, Llama 2 exhibits greater variability and a broader range of sensitivities, with a wider range of scores from approximately -15 to 20, compared to GPT-3.5’s range of -2 to 6. This suggests that Llama 2 tends to assign more extreme scores to certain verbs.




Figure 3.3: Comparison of Fightin’ Word scores for verbs between GPT-3.5 and Llama 2. The x-axis represents the scores from GPT-3.5, while the y-axis represents the scores from Llama 2. Higher scores indicate a stronger association with male protagonists, and lower scores indicate a stronger association with female protagonists.

The scatter plots present a comparison of Fightin’ Word Scores between GPT-3.5 and Llama 2 for adjectives and verbs, considering completion, contextual, instructional, and question-answer prompts. Llama 2 consistently displays more extreme scores for both adjectives and verbs across all prompt types, in contrast to GPT-3.5. Completion prompts typically yield scores clustered around zero, with Llama 2 showing a bit more variability. Contextual and question-answer prompts generate greater score variability, highlighting a higher sensitivity to context and interactivity, respectively. Instructional prompts reveal patterns akin to completion prompts but with some more extreme variations in Llama 2’s scores. In summary, these comparisons suggest that Llama 2 has a broader range of sensitivities to different prompt types, especially in contextual and question-answer formats, indicating potential influences of prompt context on the models’ evaluations.

The series of scatter plots compare the Fightin’ Word Scores between GPT-3.5 and Llama 2 for both adjectives and verbs across different genres: literary fiction, romance, and thriller. In literary fiction, most data points for both adjectives and verbs are clustered around zero for both models, indicating similar evaluations, but Llama 2 shows more extreme scores ranging from -10 to 5. In the romance genre, Llama 2 exhibits a wider range of scores from -12.5 to 7.5 for both adjectives and verbs, reflecting a higher

sensitivity to the emotional and descriptive language typical of romance stories, compared to GPT-3.5’s range of -4 to 4. While analysing the thriller genre, Llama 2 shows scores from -12.5 to 5 for both adjectives and verbs, suggesting a potential sensitivity to the intense and action-oriented language of thrillers. In comparison, GPT-3.5’s scores range from -4 to 3. Although these differences are observed, the relatively small range of scores may not meaningfully reflect distinct genre-specific language processing. This indicates that while Llama 2 might display greater variability, the overall impact of genre remains limited. Across all genres, Llama 2 consistently shows greater variability in scores compared to GPT-3.5, indicating a more nuanced language evaluation. This increased variability in Llama 2’s scores suggests a deeper sensitivity to the language used in different genres, particularly in assigning sentiment scores to both adjectives and verbs, potentially reflecting a more sophisticated understanding of the nuanced contexts within these genres.

3.5.3 Conclusion

In analysing the presence and extent of gender bias in the assignment of occupations in Dutch short stories generated using GPT-3.5 and Llama 2, several key findings emerge from the research questions and the "Fightin' Words" analysis:

Q1: user prompts influence the gender bias observed in occupation assignment. The "Fightin' Words" analysis reveals that different prompt types (completion, contextual, instructional, and question-answer) yield varying degrees of gender bias. Llama 2 consistently shows more extreme variations in scores for both adjectives and verbs across all prompt types compared to GPT-3.5, indicating a higher sensitivity to the context provided by the prompts.

Q2: the presence and extent of gender bias also vary across different literary genres. In genres like literary fiction, romance, and thriller, Llama 2 demonstrates a broader range of sensitivity to gender-associated language compared to GPT-3.5. Particularly in the romance and thriller genres, Llama 2’s scores reflect a higher sensitivity to the emotional and descriptive language typical of these genres. This suggests that genre-specific contexts amplify the gender biases inherent in the models. However, the relatively small dataset sizes limit the ability to draw definitive conclusions about genre-specific sensitivities. Further analysis with larger datasets or additional genres

is needed to better understand the impact of genre on gender bias in these models.

Q3: comparative analysis indicates that GPT-3.5 and Llama 2 differ notably in their handling of gender bias. Llama 2 tends to assign more extreme scores to adjectives and verbs associated with gender, displaying a wider range of sensitivities and potentially more pronounced gender biases. In contrast, GPT-3.5 shows a denser clustering of scores around neutral values, suggesting a more moderated response. This difference is particularly evident in the scatter plots of “Fightin’ Words” scores, where Llama 2 exhibits greater variability and sensitivity to gendered language.

4. Conclusion and Discussion

4.0.1 Summary of Results

This research explores the presence and extent of gender bias in the assignment of occupations in Dutch short stories generated using GPT-3.5 and Llama 2, specifically addressing the influence of user prompts, variations across different literary genres, and comparative differences between the two models. The findings from the analyses - Gender Distribution 3.3, Survey 3.4, and Fightin' Words 3.5 - provide insights into these aspects.

RQ: What is the presence and extent of gender bias in the assignment of occupations in Dutch short stories generated using GPT-3.5 and Llama 2?

Both GPT-3.5 and Llama 2 exhibit gender biases in the generated occupations, reflecting societal stereotypes. Male-dominated roles were more frequently assigned to male protagonists, while female protagonists were often given nurturing or less authoritative roles. GPT-3.5 tends to generate a slightly more balanced distribution of occupations, while Llama 2 shows more extreme gender biases.

Q1: How do user prompts influence the gender bias observed in occupation assignment?

User prompts significantly influenced the gender bias in occupation assignments. The type of prompt—completion, contextual, instructional, or question-answer—affected the gender distribution of generated occupations. Completion prompts often resulted in more male-dominated roles, while instructional prompts provided a more balanced distribution. Llama 2 displays greater variability and sensitivity to the prompt types, indicating that it was more influenced by the contextual and interactivity aspects of the prompts compared to GPT-3.5.

Q2: How does the presence and extent of gender bias in the assignment of occupations vary across different literary genres?

Gender bias varies across different literary genres. In literary fiction, both models generated a relatively balanced distribution of occupations. However, in the romance and thriller genres, the biases are more pronounced. Male protagonists are more frequently assigned technical and authoritative roles, while female protagonists are assigned nurturing roles. Llama 2 shows a higher sensitivity to genre-specific language, reflecting a broader range of sensitivities and potentially amplifying gender biases in certain genres.

Q3: How does GPT-3.5 differ from Llama 2 in terms of gender bias when assigning occupations to characters in generated Dutch short stories?

GPT-3.5 and Llama 2 differ notably in their handling of gender bias. GPT-3.5 generally produces a denser clustering of scores around neutral values, indicating a more moderated response. In contrast, Llama 2 assigns more extreme scores to adjectives and verbs associated with gender, displaying greater variability and sensitivity to gendered language. This suggests that Llama 2 may amplify existing gender biases more than GPT-3.5, particularly in context-specific scenarios and under the influence of certain prompt types. Given these observations, Llama 2 can be considered a benchmark for examining gender bias amplification, while GPT-3.5 generates a more balanced and less biased output.

Overall Conclusion:

The analyses highlight that both GPT-3.5 and Llama 2 exhibit gender biases reflective of societal stereotypes. User prompts and literary genres significantly influence these biases, with Llama 2 demonstrating a higher sensitivity and a broader range of gender-associated language variations. This underscores the need for balanced training data and robust bias mitigation strategies to ensure more equitable AI-generated content. Addressing these biases is crucial for developing fair and unbiased language models that can accurately represent diverse societal roles without perpetuating stereotypes.

4.1 Limitations and future work

This research faces several limitations that impact the generalisability and comprehensiveness of its findings.

Firstly, the survey includes a neutral option for gender associations, while the analysis primarily focuses on male and female genders. A significant proportion of survey responses indicate a neutral association, whereas only a small fraction of the generated occupations are assigned to neutral protagonists. This discrepancy leads to a skewed comparison and highlights the need for more inclusive research that considers different gender identities beyond the binary male-female classification.

Secondly, the survey responses reflect the perceptions of a limited and potentially non-representative sample of participants. Broader and more diverse participant pools could provide a more accurate reflection of societal gender perceptions.

Thirdly, the study primarily examines how generated occupations are assigned to genders, but it would also be valuable to investigate the reverse: how certain genders are assigned to specific occupations. For example, instead of prompting the model with “*Schrijf een Nederlandse premisse voor een kort literaire fictie verhaal over een man en beschrijf zijn beroep in een lopend verhaal.*” (“Write a Dutch premise for a short literary fiction story about a man and describe his occupation in an ongoing story.”), a reversed prompt could be used such as “*Schrijf een Nederlandse premisse voor een kort literaire fictie verhaal over een architect en beschrijf de gender van de hoofdpersoon in een lopend verhaal.*” (“Write a Dutch premise for a short literary fiction story about an architect and describe the gender of the protagonist in an ongoing story.”). Studying how certain genders are assigned to specific occupations can provide additional insights into the interplay between gender and occupations, showing potential biases from another perspective.

Fourthly, the “Fightin’ Words” technique, used to analyse gender bias in language, only counts word frequencies and does not account for the context or semantics of the words. As a result, important contextual nuances may be lost, limiting the depth and accuracy of the bias analysis. Integrating more advanced methods that incorporate context and semantic understanding can improve the accuracy of bias detection and offer deeper insights into how narratives shape gender biases.

Following from limitation four, the analysis does not extensively explore how var-

ious contextual elements within the stories, such as plot or character development, might influence the assignment of occupations to different genders. A more detailed contextual analysis could reveal additional layers of bias.

Last but not least, this research relies on datasets generated by GPT-3.5 and Llama 2, which may not fully represent the diversity of possible outputs from other LLMs. Different models might exhibit different biases, and the findings here may not be universally applicable. These models are trained on large datasets that inherently contain societal biases. While this study analyses the outputs, it does not address the root cause of these biases within the training data itself. Future research could benefit from examining and mitigating biases directly in the training datasets.

References

- Abid, A., Farooqi, M., & Zou, J. (2021). Persistent anti-muslim bias in large language models. In *Proceedings of the 2021 aaai/acm conference on ai, ethics, and society* (pp. 298–306).
- AI, M. (2023). *Introducing llama 2*. Retrieved 2024-01-15, from <https://ai.meta.com/llama/>
- Barikeri, S., Lauscher, A., Vulić, I., & Glavaš, G. (2021, August). RedditBias: A real-world resource for bias evaluation and debiasing of conversational language models. In C. Zong, F. Xia, W. Li, & R. Navigli (Eds.), *Proceedings of the 59th annual meeting of the association for computational linguistics and the 11th international joint conference on natural language processing (volume 1: Long papers)* (pp. 1941–1955). Online: Association for Computational Linguistics. Retrieved from <https://aclanthology.org/2021.acl-long.151> doi: 10.18653/v1/2021.acl-long.151
- Beutel, A. M., & Schleifer, C. (2022). Family structure, gender, and wages in stem work. *Sociological Perspectives*, 65(4), 790–819.
- Blum, L. (2004). Stereotypes and stereotyping: A moral analysis. *Philosophical papers*, 33(3), 251–289.
- Bolukbasi, T., Chang, K.-W., Zou, J. Y., Saligrama, V., & Kalai, A. T. (2016). Man is to computer programmer as woman is to homemaker? debiasing word embeddings. *Advances in neural information processing systems*, 29.
- Borchers, C., Gala, D., Gilbert, B., Oravkin, E., Bounsi, W., Asano, Y. M., & Kirk, H. (2022, July). Looking for a handsome carpenter! debiasing GPT-3 job advertisements. In C. Hardmeier, C. Basta, M. R. Costa-jussà, G. Stanovsky, & H. Gonen (Eds.), *Proceedings of the 4th workshop on gender bias in natural language processing (gebnlp)* (pp. 212–224). Seattle, Washington: Association

- for Computational Linguistics. Retrieved from <https://aclanthology.org/2022.gebnlp-1.22> doi: 10.18653/v1/2022.gebnlp-1.22
- Bosak, J., & Sczesny, S. (2011). Gender bias in leader selection? evidence from a hiring simulation study. *Sex roles*, 65, 234–242.
- Brown, T., Mann, B., Ryder, N., Subbiah, M., Kaplan, J. D., Dhariwal, P., ... others (2020). Language models are few-shot learners. *Advances in neural information processing systems*, 33, 1877–1901.
- Budig, M. J., & England, P. (2001). The wage penalty for motherhood. *American sociological review*, 204–225.
- Chang, Y., Wang, X., Wang, J., Wu, Y., Yang, L., Zhu, K., ... others (2024). A survey on evaluation of large language models. *ACM Transactions on Intelligent Systems and Technology*, 15(3), 1–45.
- Chen, B., Zhang, Z., Langrené, N., & Zhu, S. (2023). Unleashing the potential of prompt engineering in large language models: a comprehensive review. *arXiv preprint arXiv:2310.14735*.
- Cheng, M., Durmus, E., & Jurafsky, D. (2023, July). Marked personas: Using natural language prompts to measure stereotypes in language models. In A. Rogers, J. Boyd-Graber, & N. Okazaki (Eds.), *Proceedings of the 61st annual meeting of the association for computational linguistics (volume 1: Long papers)* (pp. 1504–1532). Toronto, Canada: Association for Computational Linguistics. Retrieved from <https://aclanthology.org/2023.acl-long.84> doi: 10.18653/v1/2023.acl-long.84
- Chung, J. J. Y., Kim, W., Yoo, K. M., Lee, H., Adar, E., & Chang, M. (2022). Talebrush: Sketching stories with generative pretrained language models. In *Proceedings of the 2022 chi conference on human factors in computing systems* (pp. 1–19).
- Cohen, I., Huang, Y., Chen, J., Benesty, J., Benesty, J., Chen, J., ... Cohen, I. (2009). Pearson correlation coefficient. *Noise reduction in speech processing*, 1–4.
- Cortes, P., & Pan, J. (2018). Occupation and gender. *The Oxford handbook of women and the economy*, 425–452.
- Cuddy, A. J., Wolf, E. B., Glick, P., Crotty, S., Chong, J., & Norton, M. I. (2015). Men as cultural ideals: Cultural values moderate gender stereotype content. *Journal*

- of personality and social psychology*, 109(4), 622.
- Devlin, J., Chang, M.-W., Lee, K., & Toutanova, K. (2019, June). BERT: Pre-training of deep bidirectional transformers for language understanding. In J. Burstein, C. Doran, & T. Solorio (Eds.), *Proceedings of the 2019 conference of the north American chapter of the association for computational linguistics: Human language technologies, volume 1 (long and short papers)* (pp. 4171–4186). Minneapolis, Minnesota: Association for Computational Linguistics. Retrieved from <https://aclanthology.org/N19-1423> doi: 10.18653/v1/N19-1423
- De Winter, J. C., Gosling, S. D., & Potter, J. (2016). Comparing the pearson and spearman correlation coefficients across distributions and sample sizes: A tutorial using simulations and empirical data. *Psychological methods*, 21(3), 273.
- Diekmann, A. B., & Eagly, A. H. (2000). Stereotypes as dynamic constructs: Women and men of the past, present, and future. *Personality and social psychology bulletin*, 26(10), 1171–1188.
- Dinan, E., Fan, A., Williams, A., Urbanek, J., Kiela, D., & Weston, J. (2020, November). Queens are powerful too: Mitigating gender bias in dialogue generation. In B. Webber, T. Cohn, Y. He, & Y. Liu (Eds.), *Proceedings of the 2020 conference on empirical methods in natural language processing (emnlp)* (pp. 8173–8188). Online: Association for Computational Linguistics. Retrieved from <https://aclanthology.org/2020.emnlp-main.656> doi: 10.18653/v1/2020.emnlp-main.656
- Doughman, J., Khreich, W., El Gharib, M., Wiss, M., & Berjawi, Z. (2021). Gender bias in text: Origin, taxonomy, and implications. In *Proceedings of the 3rd workshop on gender bias in natural language processing* (pp. 34–44).
- Dowling, C., & Dowling, C. (1990). *Cinderella complex*. Pocket Books New York.
- Ellemers, N. (2018). Gender stereotypes. *Annual review of psychology*, 69, 275–298.
- England, P. (2010). The gender revolution: Uneven and stalled. *Gender & society*, 24(2), 149–166.
- Erigha, M. (2015). Race, gender, hollywood: Representation in cultural production and digital media’s potential for change. *Sociology compass*, 9(1), 78–89.
- Feng, Z., Guo, D., Tang, D., Duan, N., Feng, X., Gong, M., ... Zhou, M. (2020, November). CodeBERT: A pre-trained model for programming and natural

- languages. In T. Cohn, Y. He, & Y. Liu (Eds.), *Findings of the association for computational linguistics: Emnlp 2020* (pp. 1536–1547). Online: Association for Computational Linguistics. Retrieved from <https://aclanthology.org/2020.findings-emnlp.139> doi: 10.18653/v1/2020.findings-emnlp.139
- Forsman, J. A., & Barth, J. M. (2017). The effect of occupational gender stereotypes on men’s interest in female-dominated occupations. *Sex Roles*, 76, 460–472.
- Gala, D., Khursheed, M. O., Lerner, H., O’Connor, B., & Iyyer, M. (2020, November). Analyzing gender bias within narrative tropes. In D. Bamman, D. Hovy, D. Jurgens, B. O’Connor, & S. Volkova (Eds.), *Proceedings of the fourth workshop on natural language processing and computational social science* (pp. 212–217). Online: Association for Computational Linguistics. Retrieved from <https://aclanthology.org/2020.nlpcss-1.23> doi: 10.18653/v1/2020.nlpcss-1.23
- Gerritsen, M. (2001). Changes in professional terms in the netherlands: Anglicisation and the neutralisation of gender. *Linguistics in the Netherlands*, 18(1), 101–111.
- Gira, M., Zhang, R., & Lee, K. (2022). Debiasing pre-trained language models via efficient fine-tuning. In *Proceedings of the second workshop on language technology for equality, diversity and inclusion* (pp. 59–69).
- Giray, L. (2023). Prompt engineering with chatgpt: A guide for academic writers. *Annals of Biomedical Engineering*, 1–5.
- Haleem, A., Javaid, M., & Singh, R. P. (2022). An era of chatgpt as a significant futuristic support tool: A study on features, abilities, and challenges. *BenchCouncil transactions on benchmarks, standards and evaluations*, 2(4), 100089.
- Harvie, K., Marshall-Mcaskey, J., & Johnston, L. (1998). Gender-based biases in occupational hiring decisions 1. *Journal of Applied Social Psychology*, 28(18), 1698–1711.
- Heilman, M. E. (2012). Gender stereotypes and workplace bias. *Research in organizational Behavior*, 32, 113–135.
- Hendy, A., Abdelrehim, M., Sharaf, A., Raunak, V., Gabr, M., Matsushita, H., ... Awadalla, H. H. (2023). How good are gpt models at machine translation? a comprehensive evaluation. *arXiv preprint arXiv:2302.09210*.
- Hirsh, C. E. (2009). The strength of weak enforcement: The impact of discrimination charges, legal environments, and organizational conditions on workplace

- segregation. *American Sociological Review*, 74(2), 245–271.
- Kirk, H. R., Jun, Y., Volpin, F., Iqbal, H., Benussi, E., Dreyer, F., ... Asano, Y. (2021). Bias out-of-the-box: An empirical analysis of intersectional occupational biases in popular generative language models. *Advances in neural information processing systems*, 34, 2611–2624.
- Klapwijk, P. (2022). *Kwart raakt geen boek aan, maar meerderheid leest nog steeds: dit zijn de trends van nu*. Retrieved 2024-01-15, from [https://eenvandaag.avrotros.nl/panels/opiniepanel/alle-uitslagen/item/kwart-raakt-geen-boek-aan-maar-meerderheid-leest-nog-steeds-dit-zijn-de-trends-van-nu/#:~:text=Van%20alle%20genres%20zijn%20de,35%20jaar%20\(37%20procent\)](https://eenvandaag.avrotros.nl/panels/opiniepanel/alle-uitslagen/item/kwart-raakt-geen-boek-aan-maar-meerderheid-leest-nog-steeds-dit-zijn-de-trends-van-nu/#:~:text=Van%20alle%20genres%20zijn%20de,35%20jaar%20(37%20procent)).
- Koenig, A. M., Eagly, A. H., Mitchell, A. A., & Ristikari, T. (2011). Are leader stereotypes masculine? a meta-analysis of three research paradigms. *Psychological bulletin*, 137(4), 616.
- Kurita, K., Vyas, N., Pareek, A., Black, A. W., & Tsvetkov, Y. (2019, August). Measuring bias in contextualized word representations. In M. R. Costa-jussà, C. Hardmeier, W. Radford, & K. Webster (Eds.), *Proceedings of the first workshop on gender bias in natural language processing* (pp. 166–172). Florence, Italy: Association for Computational Linguistics. Retrieved from <https://aclanthology.org/W19-3823> doi: 10.18653/v1/W19-3823
- Lauzen, M. M. (2021). The celluloid ceiling: Behind-the-scenes employment of women on the top us films of 2020. *The Center for the Study of Women in Television and Film*.
- Liu, N. F., Lin, K., Hewitt, J., Paranjape, A., Bevilacqua, M., Petroni, F., & Liang, P. (2024). Lost in the middle: How language models use long contexts. *Transactions of the Association for Computational Linguistics*, 12, 157–173.
- Liu, P., Yuan, W., Fu, J., Jiang, Z., Hayashi, H., & Neubig, G. (2023). Pre-train, prompt, and predict: A systematic survey of prompting methods in natural language processing. *ACM Computing Surveys*, 55(9), 1–35.
- Liu, R., Jia, C., Wei, J., Xu, G., Wang, L., & Vosoughi, S. (2021). Mitigating political bias in language models through reinforced calibration. In *Proceedings of the aaai conference on artificial intelligence* (Vol. 35, pp. 14857–14866).

- Lucy, L., & Bamman, D. (2021, June). Gender and representation bias in GPT-3 generated stories. In *Proceedings of the third workshop on narrative understanding* (pp. 48–55). Virtual: Association for Computational Linguistics. Retrieved from <https://aclanthology.org/2021.nuse-1.5> doi: 10.18653/v1/2021.nuse-1.5
- Marvin, G., Hellen, N., Jjingo, D., & Nakatumba-Nabende, J. (2024). Prompt engineering in large language models. In I. J. Jacob, S. Piramuthu, & P. Falkowski-Gilski (Eds.), *Data intelligence and cognitive informatics* (pp. 387–402). Singapore: Springer Nature Singapore.
- Micheltore, K., & Sassler, S. (2016). Explaining the gender wage gap in stem: does field sex composition matter? *RSF: The Russell Sage Foundation Journal of the Social Sciences*, 2(4), 194–215.
- Mirowski, P., Mathewson, K. W., Pittman, J., & Evans, R. (2023). Co-writing screenplays and theatre scripts with language models: Evaluation by industry professionals. In *Proceedings of the 2023 chi conference on human factors in computing systems* (pp. 1–34).
- Monroe, B. L., Colaresi, M. P., & Quinn, K. M. (2008). Fightin’words: Lexical feature selection and evaluation for identifying the content of political conflict. *Political Analysis*, 16(4), 372–403.
- Moss-Racusin, C. A., Dovidio, J. F., Brescoll, V. L., Graham, M. J., & Handelsman, J. (2012). Science faculty’s subtle gender biases favor male students. *Proceedings of the national academy of sciences*, 109(41), 16474–16479.
- Nadeem, M., Bethke, A., & Reddy, S. (2021, August). StereoSet: Measuring stereotypical bias in pretrained language models. In C. Zong, F. Xia, W. Li, & R. Navigli (Eds.), *Proceedings of the 59th annual meeting of the association for computational linguistics and the 11th international joint conference on natural language processing (volume 1: Long papers)* (pp. 5356–5371). Online: Association for Computational Linguistics. Retrieved from <https://aclanthology.org/2021.acl-long.416> doi: 10.18653/v1/2021.acl-long.416
- OpenAI. (2023). *Models*. Retrieved 2024-01-15, from <https://platform.openai.com/docs/models>
- Ousidhoum, N., Zhao, X., Fang, T., Song, Y., & Yeung, D.-Y. (2021). Probing toxic content in large pre-trained language models. In *Proceedings of the 59th annual*

- meeting of the association for computational linguistics and the 11th international joint conference on natural language processing (volume 1: Long papers) (pp. 4262–4274).
- Overheid, D. (2023). *Nederland bouwt eigen open taalmodel gpt-nl*. Retrieved 2023-12-22, from <https://www.digitaleoverheid.nl/nieuws/nederland-bouwt-eigen-open-taalmodel-gpt-nl/>
- Park, J. H., Shin, J., & Fung, P. (2018, October–November). Reducing gender bias in abusive language detection. In E. Riloff, D. Chiang, J. Hockenmaier, & J. Tsujii (Eds.), *Proceedings of the 2018 conference on empirical methods in natural language processing* (pp. 2799–2804). Brussels, Belgium: Association for Computational Linguistics. Retrieved from <https://aclanthology.org/D18-1302> doi: 10.18653/v1/D18-1302
- Radford, A., Wu, J., Child, R., Luan, D., Amodei, D., Sutskever, I., et al. (2019). Language models are unsupervised multitask learners. *OpenAI blog*, 1(8), 9.
- Régner, I., Thinus-Blanc, C., Netter, A., Schmader, T., & Huguet, P. (2019). Committees with implicit biases promote fewer women when they do not believe gender bias exists. *Nature human behaviour*, 3(11), 1171–1179.
- Rudinger, R., Naradowsky, J., Leonard, B., & Van Durme, B. (2018, June). Gender bias in coreference resolution. In M. Walker, H. Ji, & A. Stent (Eds.), *Proceedings of the 2018 conference of the north American chapter of the association for computational linguistics: Human language technologies, volume 2 (short papers)* (pp. 8–14). New Orleans, Louisiana: Association for Computational Linguistics. Retrieved from <https://aclanthology.org/N18-2002> doi: 10.18653/v1/N18-2002
- Sheng, E., Chang, K.-W., Natarajan, P., & Peng, N. (2019, November). The woman worked as a babysitter: On biases in language generation. In K. Inui, J. Jiang, V. Ng, & X. Wan (Eds.), *Proceedings of the 2019 conference on empirical methods in natural language processing and the 9th international joint conference on natural language processing (emnlp-ijcnlp)* (pp. 3407–3412). Hong Kong, China: Association for Computational Linguistics. Retrieved from <https://aclanthology.org/D19-1339> doi: 10.18653/v1/D19-1339
- Skorikov, V. B., & Vondracek, F. W. (2011). Occupational identity. *Handbook of*

- identity theory and research*, 693–714.
- Stanczak, K., & Augenstein, I. (2021). A survey on gender bias in natural language processing. *arXiv preprint arXiv:2112.14168*.
- Stanovsky, G., Smith, N. A., & Zettlemoyer, L. (2019, July). Evaluating gender bias in machine translation. In A. Korhonen, D. Traum, & L. Màrquez (Eds.), *Proceedings of the 57th annual meeting of the association for computational linguistics* (pp. 1679–1684). Florence, Italy: Association for Computational Linguistics. Retrieved from <https://aclanthology.org/P19-1164> doi: 10.18653/v1/P19-1164
- Thakur, V. (2023). Unveiling gender bias in terms of profession across llms: Analyzing and addressing sociological implications. *arXiv preprint arXiv:2307.09162*.
- Touvron, H., Martin, L., Stone, K., Albert, P., Almahairi, A., Babaei, Y., ... others (2023). Llama 2: Open foundation and fine-tuned chat models. *arXiv preprint arXiv:2307.09288*.
- Vanroy, B. (2023). Language resources for dutch large language modelling. *arXiv preprint arXiv:2312.12852*.
- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., ... Polosukhin, I. (2017). Attention is all you need. *Advances in neural information processing systems*, 30.
- Venkit, P. N., Srinath, M., & Wilson, S. (2022). A study of implicit bias in pretrained language models against people with disabilities. In *Proceedings of the 29th international conference on computational linguistics* (pp. 1324–1332).
- Virtudazo, A. (2024, 06). The influence of gender stereotype on the career aspirations in technology and livelihood. *International Journal of Innovative Science and Research Technology (IJISRT)*, 2247-2289. doi: 10.38124/ijisrt/IJISRT24MAY1792
- Weidinger, L., Mellor, J., Rauh, M., Griffin, C., Uesato, J., Huang, P.-S., ... others (2021). Ethical and social risks of harm from language models. *arXiv preprint arXiv:2112.04359*.
- Weidinger, L., Uesato, J., Rauh, M., Griffin, C., Huang, P.-S., Mellor, J., ... others (2022). Taxonomy of risks posed by language models. In *Proceedings of the 2022 acm conference on fairness, accountability, and transparency* (pp. 214–229).

- White, J., Fu, Q., Hays, S., Sandborn, M., Olea, C., Gilbert, H., ... Schmidt, D. C. (2023). A prompt pattern catalog to enhance prompt engineering with chatgpt. *arXiv preprint arXiv:2302.11382*.
- Wu, T., Terry, M., & Cai, C. J. (2022). Ai chains: Transparent and controllable human-ai interaction by chaining large language model prompts. In *Proceedings of the 2022 chi conference on human factors in computing systems* (pp. 1–22).
- Xu, H., Zhang, Z., Wu, L., & Wang, C.-J. (2019). The cinderella complex: Word embeddings reveal gender stereotypes in movies and books. *PloS one*, 14(11), e0225385.
- Ye, J., Chen, X., Xu, N., Zu, C., Shao, Z., Liu, S., ... others (2023). A comprehensive capability analysis of gpt-3 and gpt-3.5 series models. *arXiv preprint arXiv:2303.10420*.
- Yuan, A., Coenen, A., Reif, E., & Ippolito, D. (2022). Wordcraft: story writing with large language models. In *27th international conference on intelligent user interfaces* (pp. 841–852).
- Zhang, Y., Sun, S., Galley, M., Chen, Y.-C., Brockett, C., Gao, X., ... Dolan, B. (2020, July). DIALOGPT : Large-scale generative pre-training for conversational response generation. In A. Celikyilmaz & T.-H. Wen (Eds.), *Proceedings of the 58th annual meeting of the association for computational linguistics: System demonstrations* (pp. 270–278). Online: Association for Computational Linguistics. Retrieved from <https://aclanthology.org/2020.acl-demos.30> doi: 10.18653/v1/2020.acl-demos.30
- Zhao, J., Wang, T., Yatskar, M., Ordonez, V., & Chang, K.-W. (2018, June). Gender bias in coreference resolution: Evaluation and debiasing methods. In M. Walker, H. Ji, & A. Stent (Eds.), *Proceedings of the 2018 conference of the north American chapter of the association for computational linguistics: Human language technologies, volume 2 (short papers)* (pp. 15–20). New Orleans, Louisiana: Association for Computational Linguistics. Retrieved from <https://aclanthology.org/N18-2003> doi: 10.18653/v1/N18-2003
- Zhao, J., Zhou, Y., Li, Z., Wang, W., & Chang, K.-W. (2018, October-November). Learning gender-neutral word embeddings. In E. Riloff, D. Chiang, J. Hockenmaier, & J. Tsujii (Eds.), *Proceedings of the 2018 conference on empiri-*

cal methods in natural language processing (pp. 4847–4853). Brussels, Belgium: Association for Computational Linguistics. Retrieved from <https://aclanthology.org/D18-1521> doi: 10.18653/v1/D18-1521

Zhao, Z., Wallace, E., Feng, S., Klein, D., & Singh, S. (2021). Calibrate before use: Improving few-shot performance of language models. In *International conference on machine learning* (pp. 12697–12706).

A. Gender Distribution per Prompt

The following tables A.1 and A.2 present the top 5 most common generated occupations for each prompt type in GPT-3.5 and Llama 2.

A.1 Gender Distribution per Prompt Type GPT-3.5

Prompt Type	Occupation	# Men	# Women	# Neutral	Total
Completion	<i>detective</i>	218	81	4	303
Completion	<i>chef</i>	200	68	0	268
Completion	<i>advocaat</i>	20	64	0	84
Completion	<i>archeoloog</i>	42	23	0	65
Completion	<i>piloot</i>	35	11	0	46
Contextual	<i>architect</i>	105	36	0	141
Contextual	<i>bibliothecaris</i>	69	45	0	114
Contextual	<i>advocaat</i>	50	58	0	108
Contextual	<i>chef</i>	65	30	0	95
Contextual	<i>forensisch psycholoog</i>	34	38	0	72
Instructional	<i>architect</i>	128	137	1	266
Instructional	<i>restaurateur</i>	32	87	1	120
Instructional	<i>straatmuzikant</i>	62	13	0	75
Instructional	<i>archeoloog</i>	38	32	5	75
Instructional	<i>forensisch psycholoog</i>	7	47	0	54

Prompt Type	Occupation	# Men	# Women	# Neutral	Total
Question-answer	<i>architect</i>	93	41	1	135
Question-answer	<i>chef</i>	77	26	1	100
Question-answer	<i>advocaat</i>	19	49	0	68
Question-answer	<i>bibliothecaris</i>	35	27	2	64
Question-answer	<i>forensisch psycholoog</i>	26	35	1	62

A.2 Gender Distribution per Prompt Type Llama

2

Prompt Type	Occupation	# Men	# Women	# Neutral	Total
Completion	<i>detective</i>	180	158	9	347
Completion	<i>schilder</i>	187	87	0	274
Completion	<i>schrijver</i>	153	63	0	216
Completion	<i>astronaut</i>	34	61	0	95
Completion	<i>fotograaf</i>	39	1	0	40
Contextual	<i>journalist</i>	42	107	3	152
Contextual	<i>architect</i>	81	65	0	146
Contextual	<i>ambtenaar</i>	1	81	0	82
Contextual	<i>timmerman</i>	67	0	0	67
Contextual	<i>beeldhouwer</i>	62	0	0	62
Instructional	<i>architect</i>	87	240	3	330
Instructional	<i>fotograaf</i>	32	105	0	137
Instructional	<i>journalist</i>	8	96	6	110
Instructional	<i>detective</i>	80	22	5	107
Instructional	<i>schilder</i>	39	60	0	99
Question-Answer	<i>architect</i>	100	55	3	158
Question-Answer	<i>zakenpersoon</i>	113	9	0	122
Question-Answer	<i>journalist</i>	5	107	1	113
Question-Answer	<i>schilder</i>	73	19	2	94
Question-Answer	<i>kunstenaar</i>	5	63	3	71

B. Gender Distribution per Genre

The following tables ?? and ?? present the top 10 most common generated occupations for each genre in GPT-3.5 and Llama 2.

B.1 Gender Distribution per Genre GPT-3.5

Genre	Occupation	# Men	# Women	# Neutral	Total
Literary fiction	<i>architect</i>	98	97	1	196
Literary fiction	<i>bibliothecaris</i>	87	50	1	138
Literary fiction	<i>archeoloog</i>	64	56	2	122
Literary fiction	<i>restaurateur</i>	12	83	0	95
Literary fiction	<i>advocaat</i>	30	62	0	92
Literary fiction	<i>detective</i>	49	19	1	69
Literary fiction	<i>chef</i>	40	15	0	55
Literary fiction	<i>straatmuzikant</i>	36	4	0	40
Literary fiction	<i>kunstenaar</i>	13	10	0	23
Literary fiction	<i>vuurtorenwachter</i>	16	4	0	20
Romance	<i>chef</i>	294	111	1	406
Romance	<i>architect</i>	171	111	1	283
Romance	<i>bloemist</i>	12	42	1	55
Romance	<i>fotograaf</i>	19	24	2	45
Romance	<i>straatmuzikant</i>	33	11	0	44
Romance	<i>restaurateur</i>	8	33	0	41
Romance	<i>bibliothecaris</i>	17	22	1	40
Romance	<i>advocaat</i>	4	28	0	32
Romance	<i>piloot</i>	24	7	0	31

Genre	Occupation	# Men	# Women	# Neutral	Total
Romance	<i>botanicus</i>	5	19	0	24
Thriller	<i>detective</i>	216	83	3	302
Thriller	<i>forensisch psycholoog</i>	64	111	1	176
Thriller	<i>advocaat</i>	55	83	0	138
Thriller	<i>architect</i>	78	16	0	94
Thriller	<i>forensisch accountant</i>	57	19	0	76
Thriller	<i>archeoloog</i>	36	25	4	65
Thriller	<i>forensisch patholoog</i>	23	35	1	59
Thriller	<i>patholoog</i>	29	14	1	44
Thriller	<i>fotograaf</i>	21	12	0	33
Thriller	<i>forensisch onderzoeker</i>	14	17	0	31

B.2 Gender Distribution per Genre Llama 2

Genre	Occupation	# Men	# Women	# Neutral	Total
Literary fiction	<i>schilder</i>	177	80	2	259
Literary fiction	<i>schrijver</i>	154	53	2	209
Literary fiction	<i>architect</i>	4	134	0	138
Literary fiction	<i>astronaut</i>	33	60	0	93
Literary fiction	<i>ambtenaar</i>	6	83	0	89
Literary fiction	<i>kunstenaar</i>	14	70	4	88
Literary fiction	<i>timmerman</i>	73	0	0	73
Literary fiction	<i>geluidstechnicus</i>	0	66	0	66
Literary fiction	<i>fotograaf</i>	6	35	0	41
Literary fiction	<i>meubelmaker</i>	25	0	0	25
Romance	<i>architect</i>	251	160	1	412
Romance	<i>schilder</i>	138	89	0	227
Romance	<i>fotograaf</i>	83	118	0	201
Romance	<i>schrijver</i>	40	29	0	69
Romance	<i>beeldhouwer</i>	61	0	0	61
Romance	<i>eventplanner</i>	0	42	3	45

Genre	Occupation	# Men	# Women	# Neutral	Total
Romance	<i>componist</i>	0	40	0	40
Romance	<i>chef</i>	5	11	13	29
Romance	<i>zakenman</i>	18	1	0	19
Romance	<i>journalist</i>	3	11	1	15
Thriller	<i>detective</i>	301	200	21	522
Thriller	<i>journalist</i>	49	282	9	340
Thriller	<i>zakenman</i>	110	11	0	121
Thriller	<i>architect</i>	14	65	5	84
Thriller	<i>ondernemer</i>	14	40	0	54
Thriller	<i>advocaat</i>	38	10	1	49
Thriller	<i>detective</i>	17	0	2	19
Thriller	<i>huurmoordenaar</i>	8	4	0	12
Thriller	<i>onderzoeksjournalist</i>	0	9	2	11
Thriller	<i>ingenieur</i>	10	0	0	10

C. Survey

C.1 Occupations

The survey occupations with their English translations. An asterisk (*) indicates that the occupation appears in the GPT-3.5 generated occupations. Likewise, a dagger (†) indicates that the occupation appears in the Llama 2 generated occupations.

Occupation

onderwijsassistent (teaching assistant)

software developer

sportinstructeur (sports instructor)

*conservator (curator)**

beeldhouwer (sculptor)†*

choreograaf (choreographer)†

verkoper (seller)†

privédetective (private detective)

huurmoordenaar (assassin)†*

beleidsmedewerker (policy officer)

adviseur (advisor)†

inkoper (buyer)

natuurkundige (physicist)

CIA-agent

grafische vormgever (graphic designer)

analist (analyst)

cartograaf (cartographer)†

Continued on next page

Table C.1 – continued from previous page

Occupation
<i>televisieproducent (television producer)</i>
<i>verkoopmedewerker (sales associate)</i>
<i>museumbeheerder (museum manager)</i>
<i>human source management</i>
<i>astroloog (astroloog)[†]</i>
<i>boekhouder (bookkeeper)*</i>
<i>model</i>
<i>theateracteur (theatre actor)</i>
<i>web designer</i>
<i>therapeut (therapist)*</i>
<i>redacteur (editor)[†]</i>
<i>kinderopvang medewerker (childcare worker)</i>
<i>public relations</i>
<i>verpleegkundige (nurse)*</i>
<i>journalist^{*†}</i>
<i>receptionist</i>
<i>software engineer</i>
<i>docent (teacher)^{*†}</i>
<i>forensisch psycholoog (forensic psychologist)*</i>
<i>kunstenaar (artist)^{*†}</i>
<i>architect^{*†}</i>
<i>evenementenplanner (event planner)[†]</i>
<i>kassamedewerker (cashier)</i>
<i>financieel adviseur (financial advisor)</i>
<i>postbezorger (mail deliverer)</i>
<i>boer (farmer)^{*†}</i>
<i>office manager</i>
<i>omroeper (announcer)</i>
<i>fotograaf (photographer)^{*†}</i>

Continued on next page

Table C.1 – continued from previous page

Occupation
<i>restaurator (restorer)*</i>
<i>forensisch onderzoeker (forensic researcher)*†</i>
<i>professor*</i>
<i>product ontwerper (product designer)</i>
<i>data-analist (data analyst)*†</i>
<i>game designer</i>
<i>notaris (notary)</i>
<i>schrijver (writer)*†</i>
<i>meubelontwerper (furniture designer)</i>
<i>musicus (musician)*</i>
<i>teamleider (team leader)</i>
<i>literatuuronderzoeker (literature researcher)</i>
<i>piloot (pilot)*</i>
<i>auteur (author)†</i>
<i>schoonmaker (cleaner)</i>
<i>componist (composer)†</i>
<i>programmeur (programmer)</i>
<i>kunstgaleriehouder (art gallery holder)</i>
<i>militair (soldier)</i>
<i>chef†</i>
<i>monteur (mechanic)</i>
<i>advocaat (lawyer)*†</i>
<i>archivaris (archivist)*</i>
<i>taalkundige (linguist)</i>
<i>zanger (singer)</i>
<i>accountant*†</i>
<i>belastingadviseur (tax advisor)</i>
<i>fiscalist (tax specialist)</i>
<i>interieurontwerper (interior designer)*†</i>

Continued on next page

Table C.1 – continued from previous page

Occupation
<i>regisseur (director)</i>
<i>danser (dancer)[†]</i>
<i>tatoeëerder (tattoo artist)</i>
<i>ambulance hulpverlener (paramedic)</i>
<i>archiefbbeheerder (archive manager)</i>
<i>stenograaf (stenographer)</i>
<i>jurist (lawyer)</i>
<i>wiskundige (mathematician)</i>
<i>notulist (court reporter)</i>
<i>arts (doctor)^{*†}</i>
<i>bloemist (florist)^{*†}</i>
<i>museummedewerker (museum employee)</i>
<i>acteur (actor)</i>
<i>tolk (interpreter)</i>
<i>ingenieur (engineer)^{*†}</i>
<i>schilder (painter)^{*†}</i>
<i>bibliothecaris (librarian)^{*†}</i>
<i>zakenpersoon (business person)^{*†}</i>
<i>archeoloog (archeologist)^{*†}</i>
<i>industrieel ontwerper (industrial designer)</i>
<i>administrateur (administrator)</i>
<i>vertaler (translator)[†]</i>
<i>risicomanager (risk manager)</i>
<i>marketeer^{*†}</i>
<i>filmacteur (movie actor)</i>
<i>illustrator^{*†}</i>
<i>beleidsadviseur (policy advisor)</i>
<i>psycholoog (psychologist)^{*†}</i>

C.2 Pearson and Spearman Correlation Coefficients Plots

Figures C.1 and C.2 show the correlation coefficients for GPT-3.5 and Llama 2 respectively, between the percentage distribution of the model's generated data and survey responses for male and female associations, thus the neutral associations are left out. The Pearson correlation coefficients are displayed in the top two plots, while the Spearman correlation coefficients are in the bottom two plots. Each plot provides a regression or monotonic line to visualise the trend.

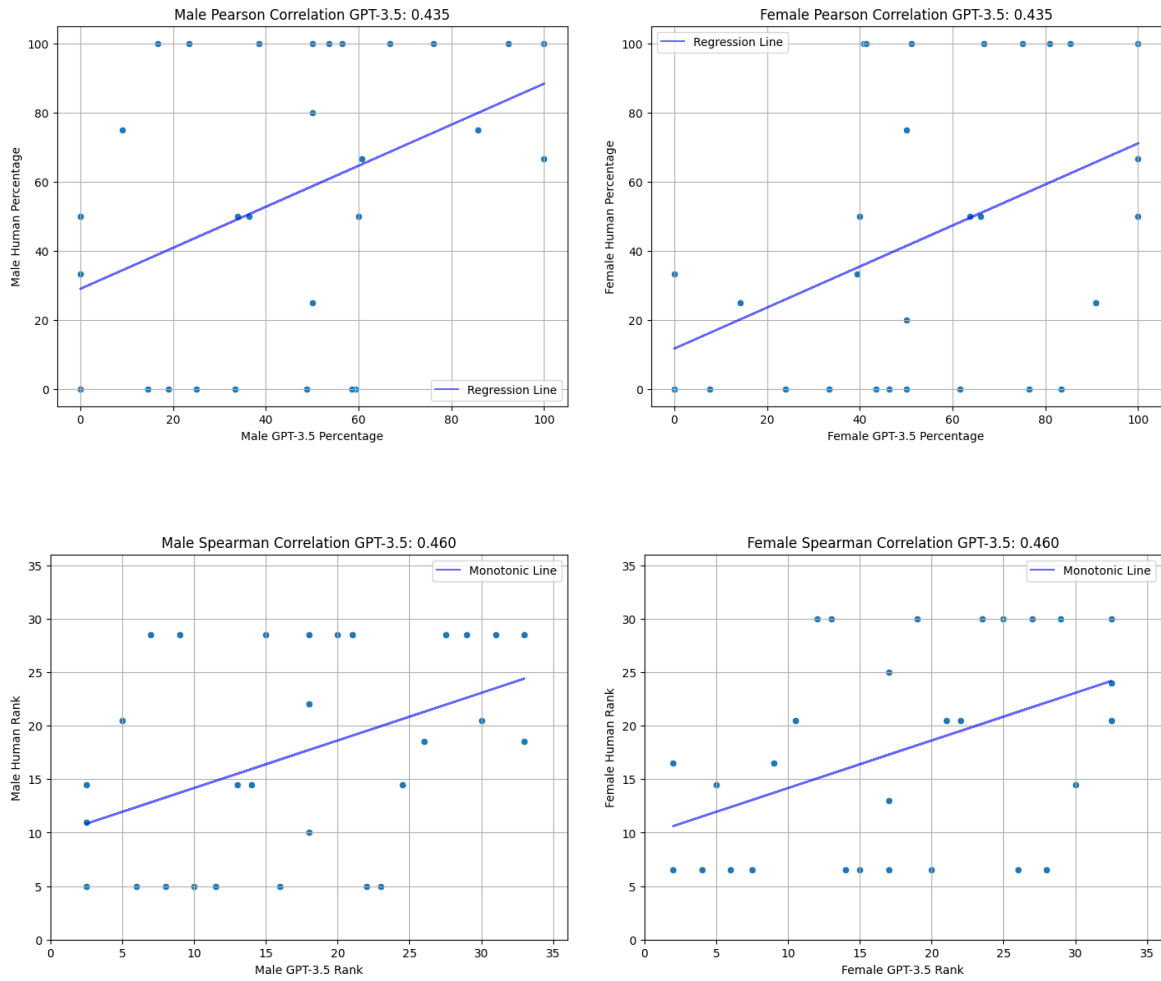


Figure C.1: Plots of the Pearson and Spearman correlation coefficients between GPT-3.5 and survey responses.

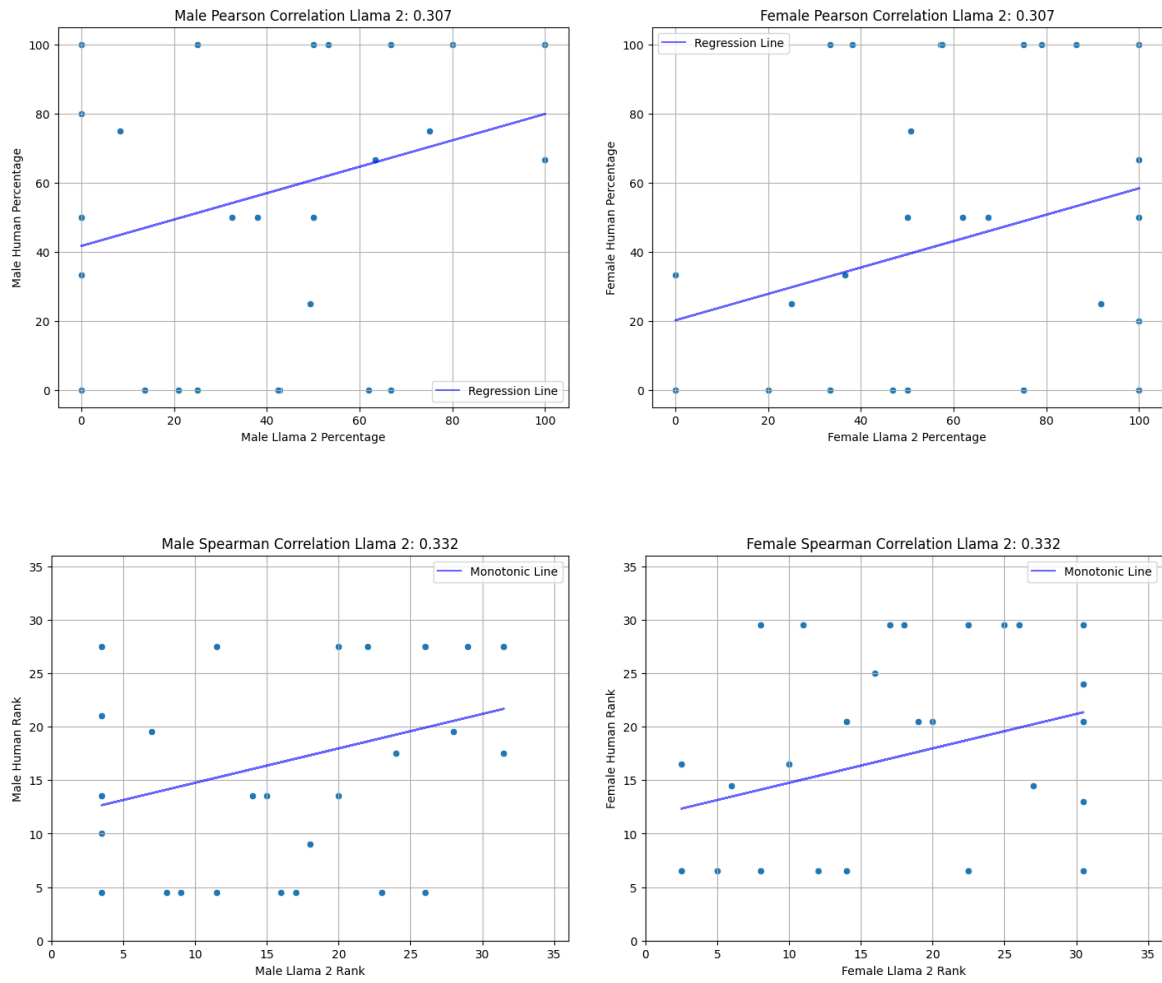


Figure C.2: Plots of the Pearson and Spearman correlation coefficients between Llama 2 and survey responses.

D. Fightin' Words

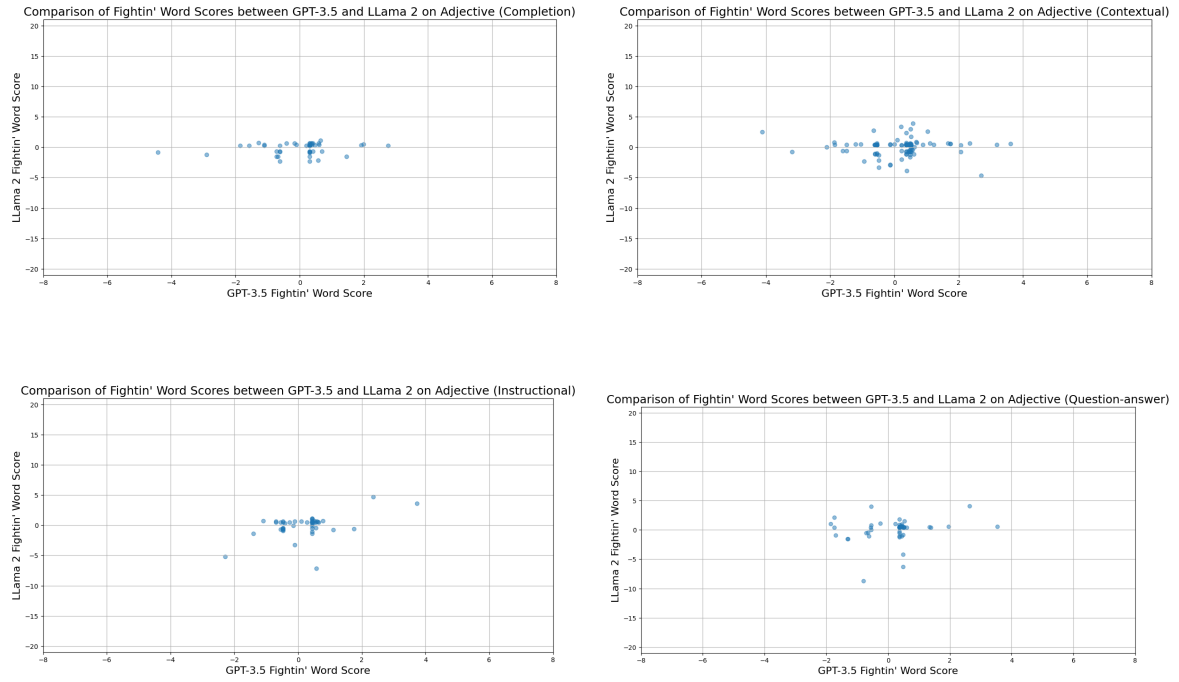


Figure D.1: Comparison of Fightin' Word scores between GPT-3.5 and Llama 2 on adjectives for different prompt types: (a) Completion, (b) Contextual, (c) Instructional, and (d) Question-answer.

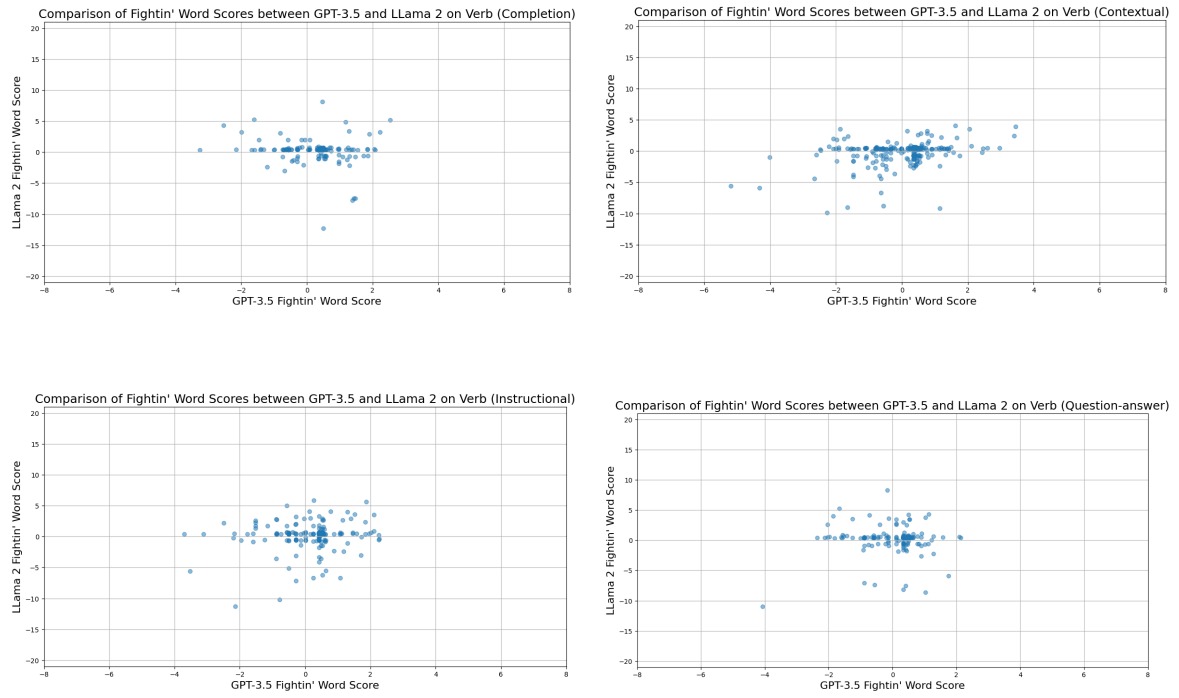


Figure D.2: Comparison of Fightin' Word scores between GPT-3.5 and Llama 2 on verbs for different prompt types: (a) Completion, (b) Contextual, (c) Instructional, and (d) Question-answer.




Figure D.3: Comparison of Fightin' Word scores between GPT-3.5 and Llama 2 on adjectives for genres prompt types: (a) Literary Fiction, (b) Romance, (c) Thriller.




Figure D.4: Comparison of Fightin' Word scores between GPT-3.5 and Llama 2 on verbs for genres prompt types: (a) Literary Fiction, (b) Romance, (c) Thriller.